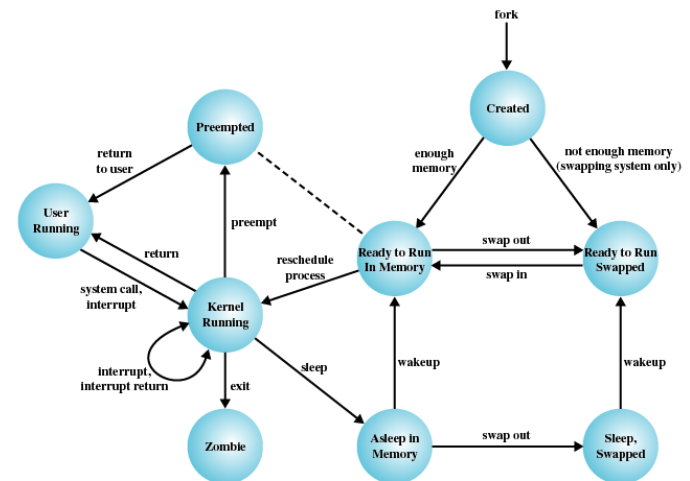




Operating Systems & Computer Networks

Introduction and Motivation



1. Introduction and Motivation

2. Subsystems, Interrupts and System Calls

3. Processes

4. Memory

5. Scheduling

6. I/O and File System

7. Booting, Services, and Security

Motivation



Finder File Edit View Go Window Help

Ubuntu2 [Running]

Text Editor

guenes@guenes-VirtualBox:~

1

VIM - Vi IMproved

version 7.3.429

by Bram Moolenaar et al.

Modified by pkg-vim-maintainers@lists.ubuntu.com

Vim is open source and freely distributable

Become a registered Vim user!

type :help register<Enter> for information

*Untitled Document 3 - gedit

Open Save Undo

Source Code

*Untitled Document 3

1 This is a text document.

Plain Text Tab Width: 4 Ln 1, Col 25

DES-Testbed | Making exp

des-testbed.net

Apps Audiobook Gündung Haus iPad/iPhone Jobs Smart-TV Softwareprojekte Go

Content management Site building Site configuration Project administration User management Reports Rules Advanced help Help 7/1 Log out guenes

Freie Universität Berlin

Home DES-Testbed DES-Workgroup DES-Projects Publications Student Research Seminar Open these topics Intern

Telematics Computer Systems

DES-Testbed

- Home
- DES-Testbed
- DES-Workgroup
 - Members
 - Projects
 - Jobs
- DES-Projects
- Publications
- Student Research Seminar

The DES-Testbed

The Distributed Embedded System (DES) research group studies real world wireless multi-hop networks (WMHN), the performance and applicability of abstract algorithms, as well as the common assumptions of simulations and analytical methods by testbed-based research.

The DES-Testbed is a hybrid wireless multi-transceiver testbed for long-term studies. It consists of a wireless mesh network (WMN) and a wireless sensor network (WSN). Our hybrid devices are called **DES-Nodes**. Currently, the testbed consists of 115 wireless mesh routers equipped with three or more IEEE 802.11a/b/g network adapters and the same number of wireless sensor nodes of type MSB-A2. In addition we have a virtualizer running several virtual machines (VMs) that recreate the testbed's topology and its

Example visualization of one part of the DES-Testbed

To view the current state of the testbed use our DES-Vis application.

Short News

- Just got th paper got i years BWO Details at l 31 weeks 4
- Great start accepted a full paper i at http://i- 34 weeks 6
- Don't miss workshop i IOC 2013 i successor i NovaElev

OSCN_01_Introduction_and_Motivation.pptx

In der Präsentation suchen

Start Designs Tabellen Diagramme SmartArt Übergänge Animationen Bildschirmpräsentation Überprüfen >> ^>

Folien

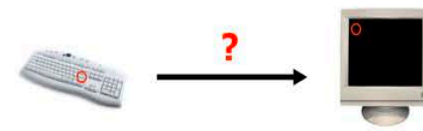
Schriftart Absatz Einfügen Format

Neue Folie

Operating System (OS) Example

Freie Universität Berlin

- What happens if one presses a key on the computer?



TI 3: Operating Systems and Computer Networks

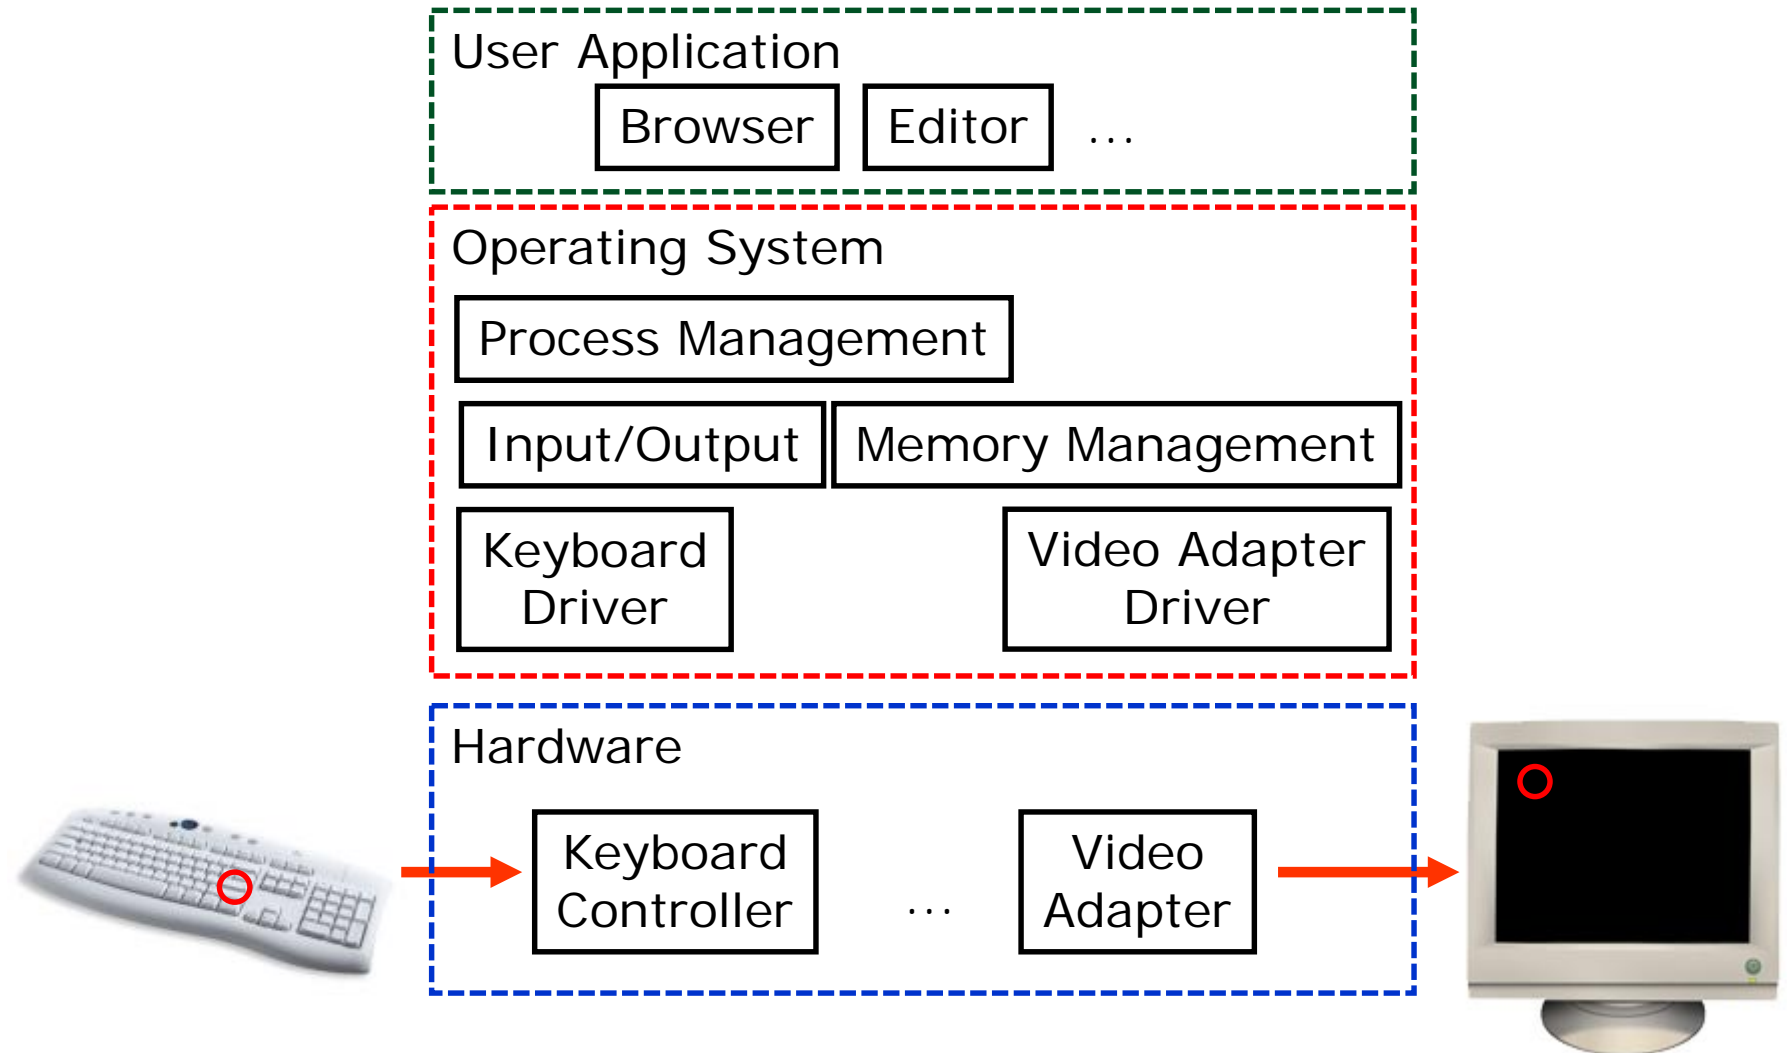
1.4

Folie 4 von 28 78%

- What happens if one presses a key on the computer?



Operating System Example



User Interface
(Shell, GUI, ...)

User Applications


System Interface
(system calls, C functions)

Operating System / Kernel

Hardware Interface
(ISA, I/O Ports, ...)

Hardware

- System interface is the only way for user applications to interact with the operating system.
- System interface consists of system calls (supervisor calls) → POSIX.



- High-level programming languages hide systems calls in library routines.

- Portable Operating System Interface (POSIX)
 - <http://standards.ieee.org/develop/wg/POSIX.html>
- POSIX defines
 - Application programming interface (API)
 - Command line shells
 - Utilities
- UNIX like Operating Systems
- POSIX oriented operating systems
 - Unix
 - Linux
 - Windows
 - Mac OS X
 - ...

Tasks of an Operating System

- Typical services of a **general** purpose OS includes:
 - Program execution
 - Access to I/O-devices
 - Hardware abstraction
 - Controlled access to files
 - Non-volatile memory
 - Access control
 - Security / user management
 - Error detection and error handling
 - Both hardware and software
 - Logging
- Special purpose operating systems focus on different services, e.g., real-time or communication requirements.

Goals of an Operating System

- Ease of use for users and programmers
- Efficiency when managing limited resources
- Possibility to evolve
 - New hardware standards
 - Changing user requirements

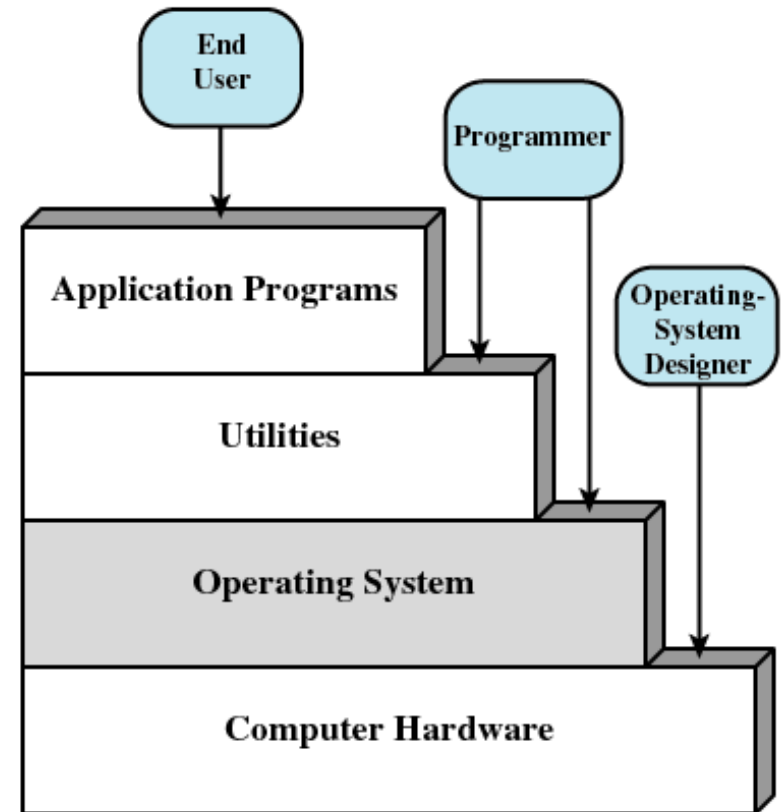
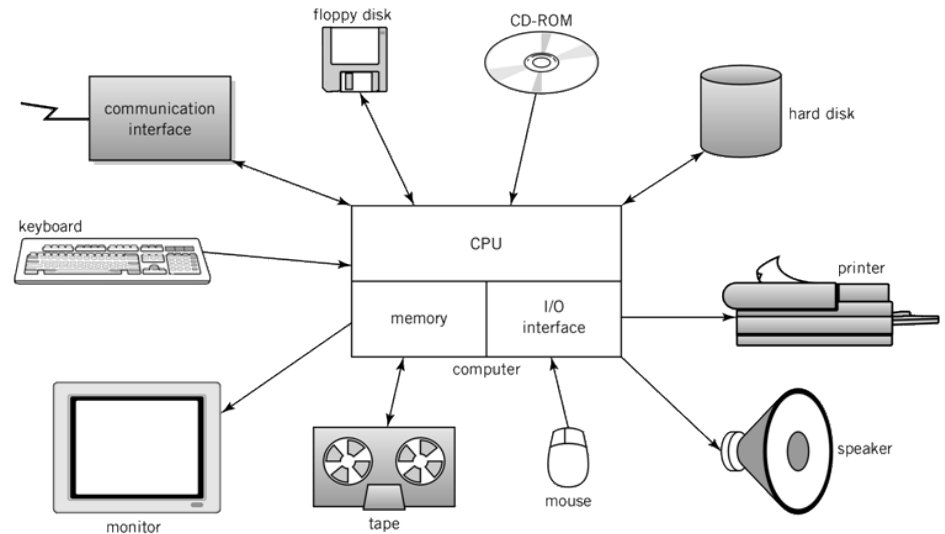


Figure 2.1 Layers and Views of a Computer System

- Hardware provides the basic computing resources such as
 - Processor(s)
 - Memory
 - Persistent storage
 - Network connection



Englander: The Architecture of Computer
Hardware and Systems Software, 2nd edition
Chapter 1, Figure 01-06

- OS **virtualizes** resources to permit controlled sharing and isolation
 - virtual instances of a resource are created
- OS provides virtual resources for user applications

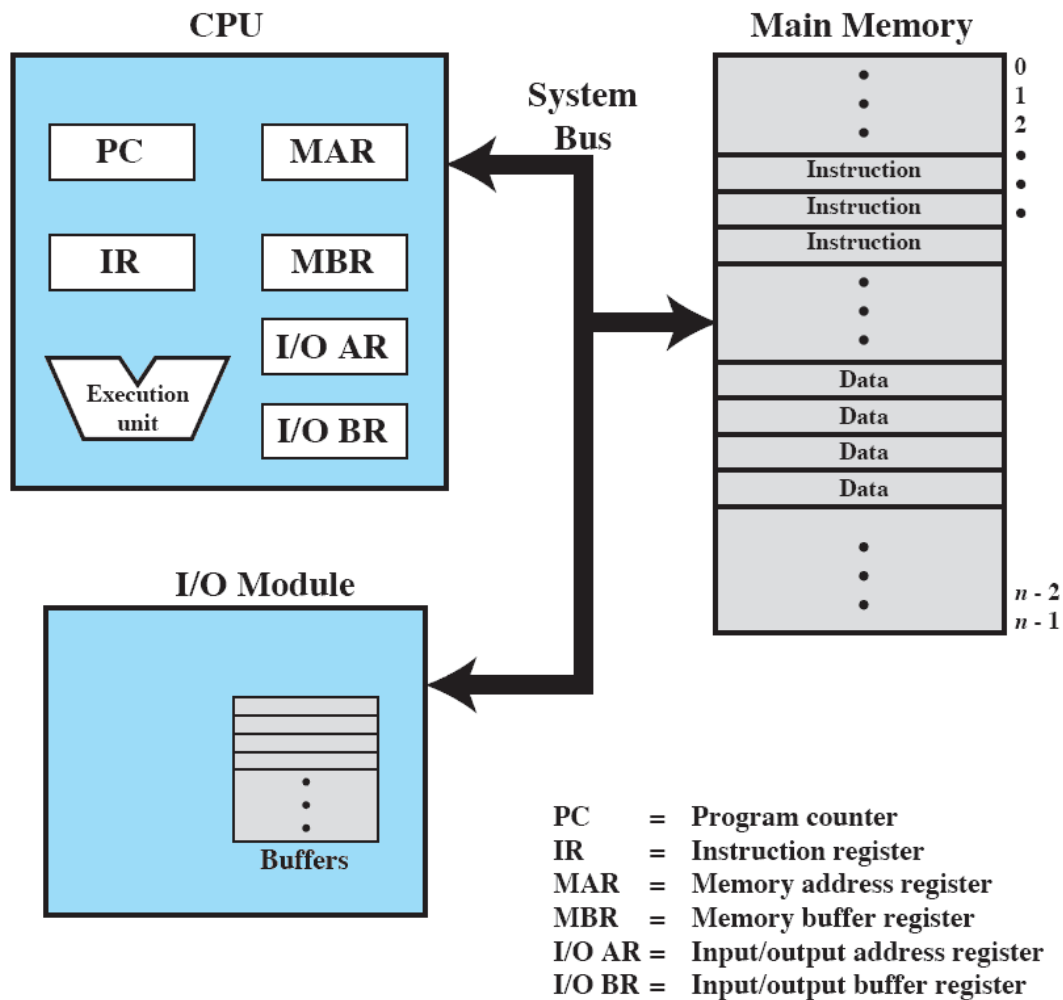



Figure 1.1 Computer Components: Top-Level View

- Virtual resources and corresponding real resources:
 - Processes processor(s)
 - Virtual Memory main memory
 - Files persistent memory
 - Ports network adapter
- Advantages:
 - Easy to use through procedural interface (system calls)
 - Secure against hardware and software errors or manipulation

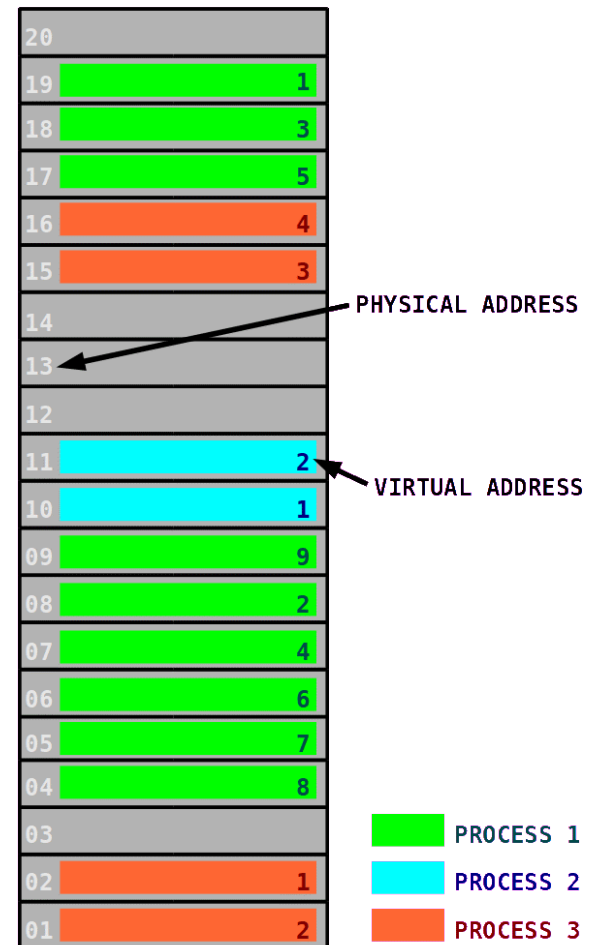
- Number of processes is not limited by the number of processors:

Multitasking

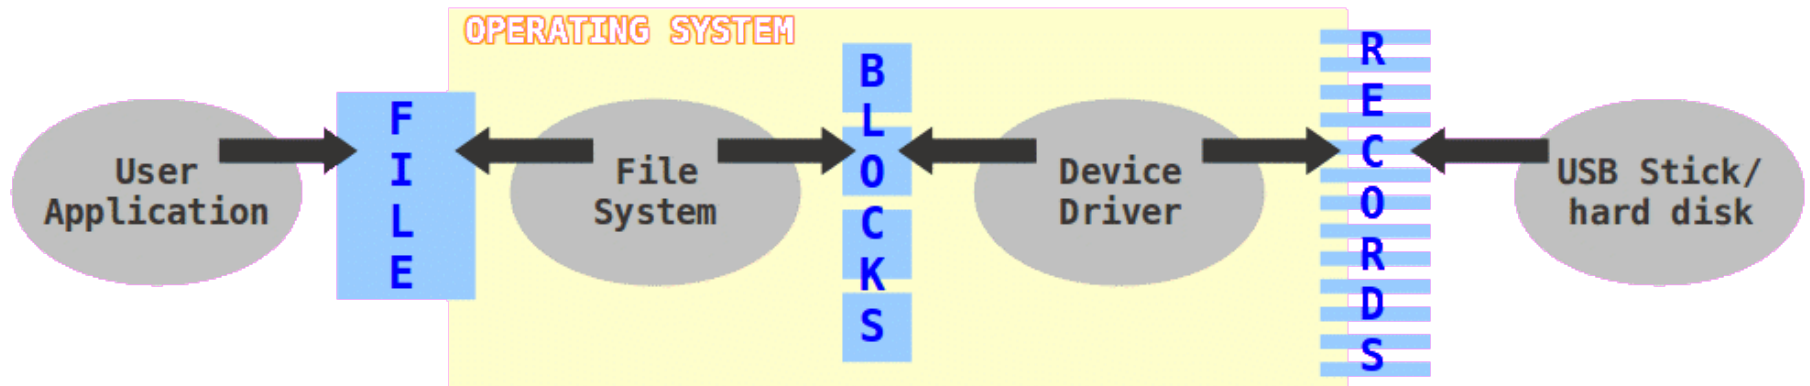
- Processor is used efficiently:
Time is not wasted by processes that are waiting on I/O devices
- Reduced latency (=response time)
- Different **process states**, e.g.,
 - running – executing
 - pending – ready to execute
 - blocked – not ready to execute



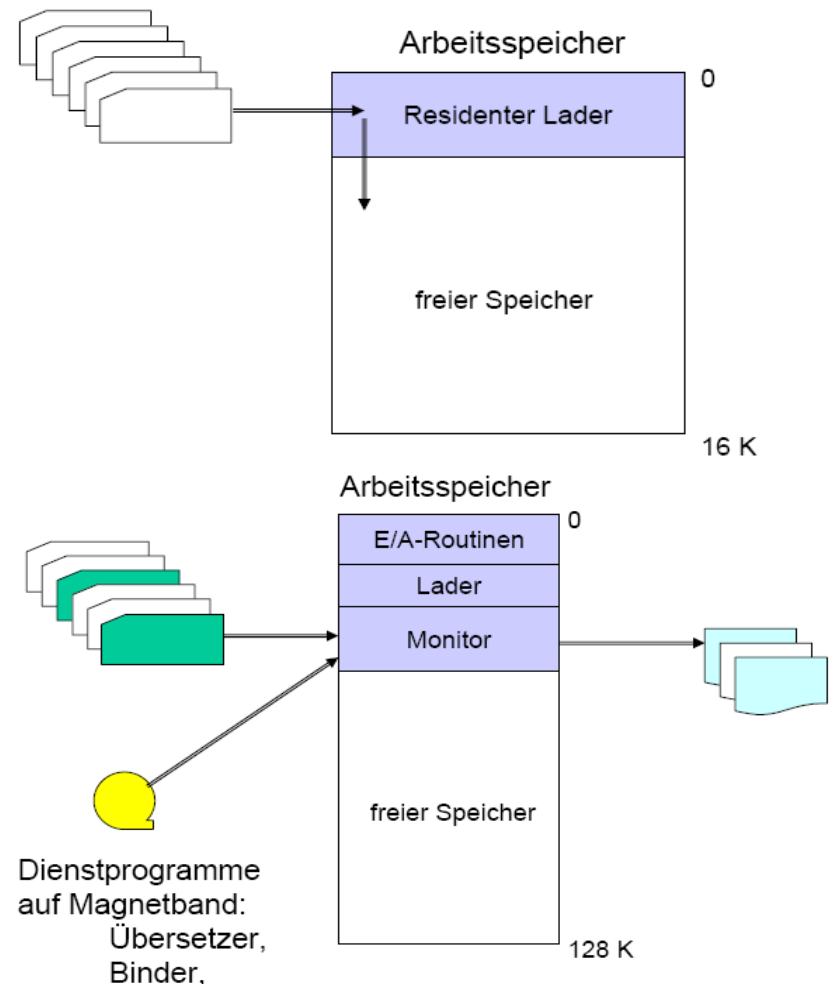
- Managed by the Memory Management Unit (MMU)
- Transportability:
 - position independent code – program does not depend on memory architecture
- Security:
 - memory access is restricted to memory units “owned” by a process
- Efficiency:
 - external fragmentation is avoided



- Managed by a file system
- Persistent objects for long-term data storage
- Stored in secondary memory (e.g., tape, hard disk, USB flash drive)
- Similar to virtual memory - file name instead of virtual address

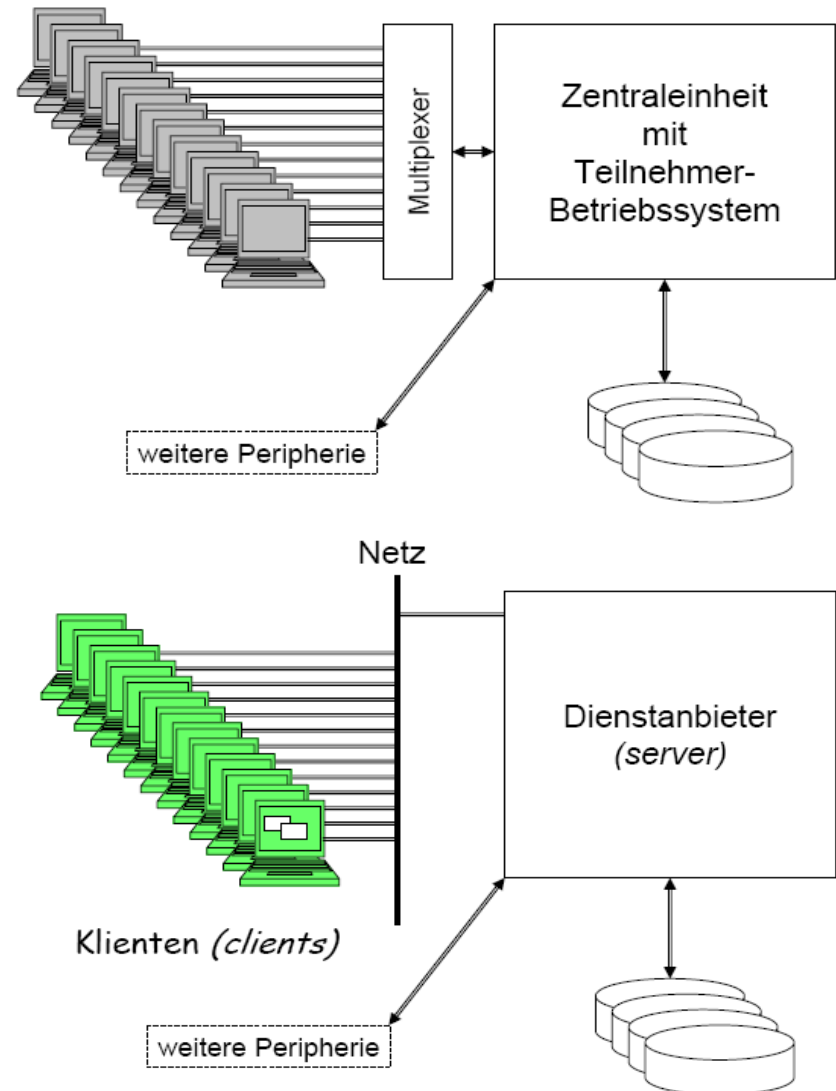


- Development of operating systems follows changes in computer architecture
- Loader (1950, IBM 704)
 - Loads programs into memory
- Batch System (1960, IBM 7090, Zuse Z 23, Telefunken TR4)
 - Processing of jobs stored on punch cards
 - Manual job control by human operator

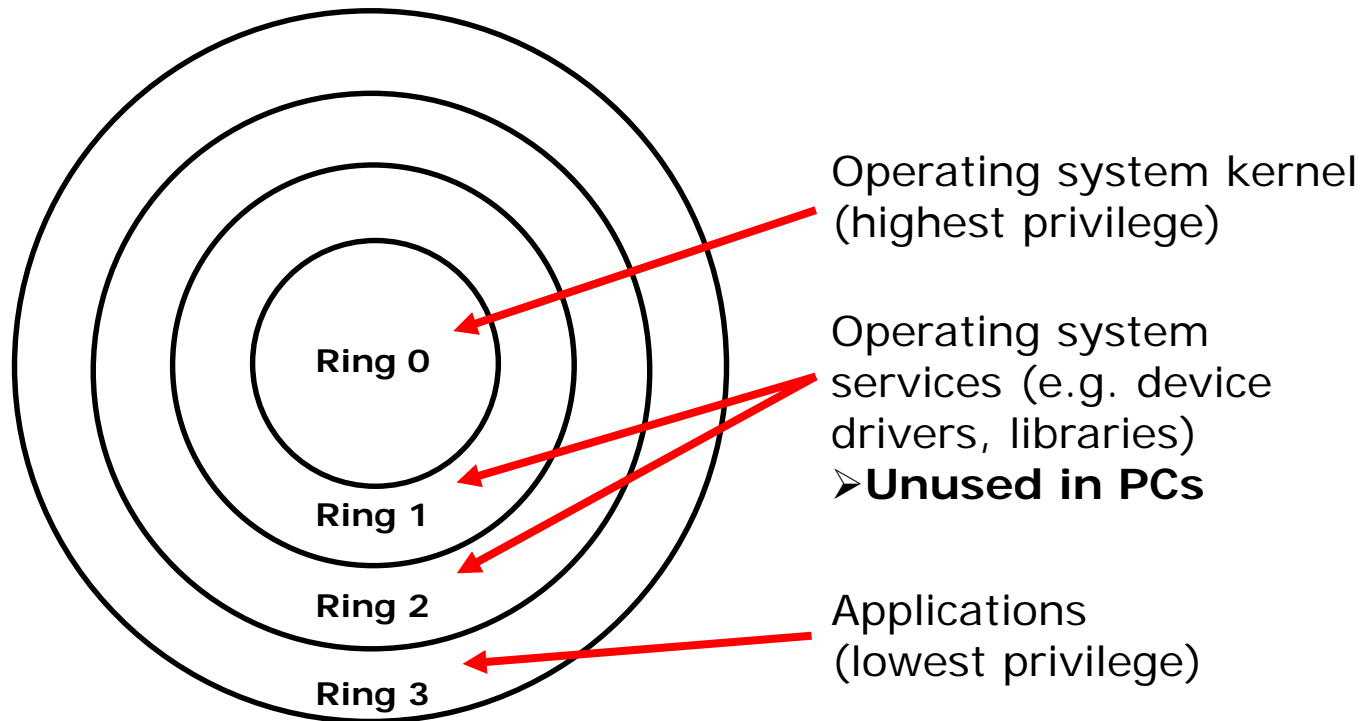


History of Operating Systems

- Multi-User / Time Sharing Systems
(1970, IBM OS/360, TSS, T.H.E., Multics, UNIX)
 - Many terminals connected to one computer
 - Interactive control for users
 - Multitasking
- Personal Computing und Client/Server
(1980/90, Apple Lisa, MS Windows, Linux, Solaris, HP-UX)
 - Intelligent workstations
 - GUI / Window mode




- Hardware provides hierarchical privilege levels
 - Inner rings have access to outer rings' resources
 - Outer rings may access inner rings through predefined gateways




- Kernel implements basic layer of abstraction
- Runs with full access to hardware (Ring 0)
- Context Switch: switching from one process to another
 - A certain amount of time is required for doing the administration, e.g., saving and loading registers.

- Monolithic Kernel

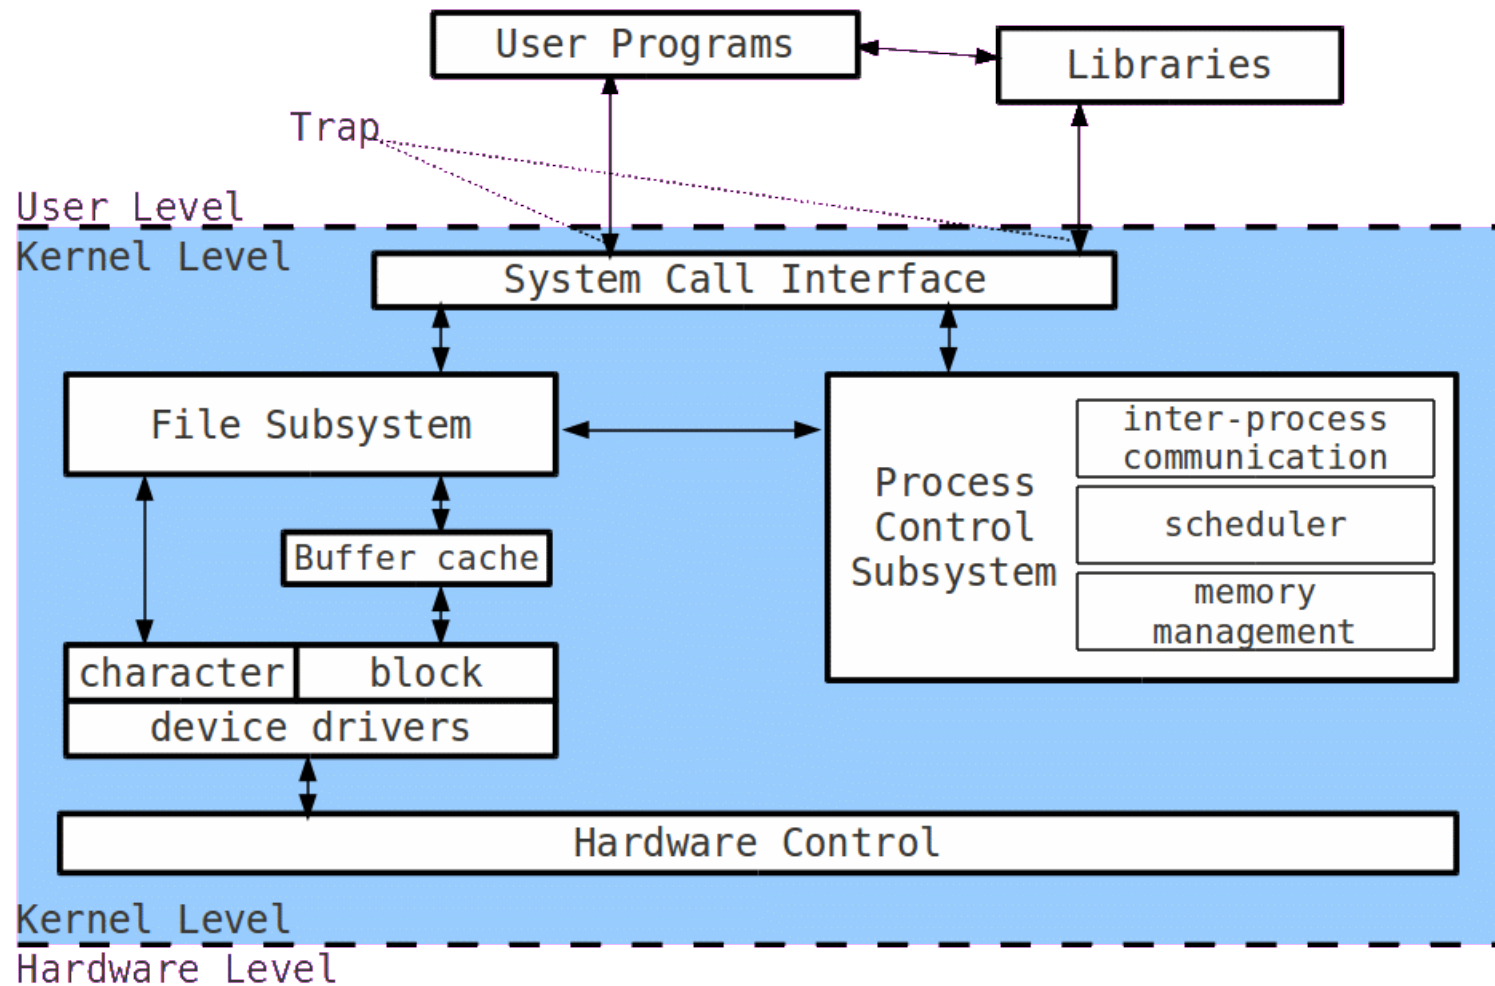


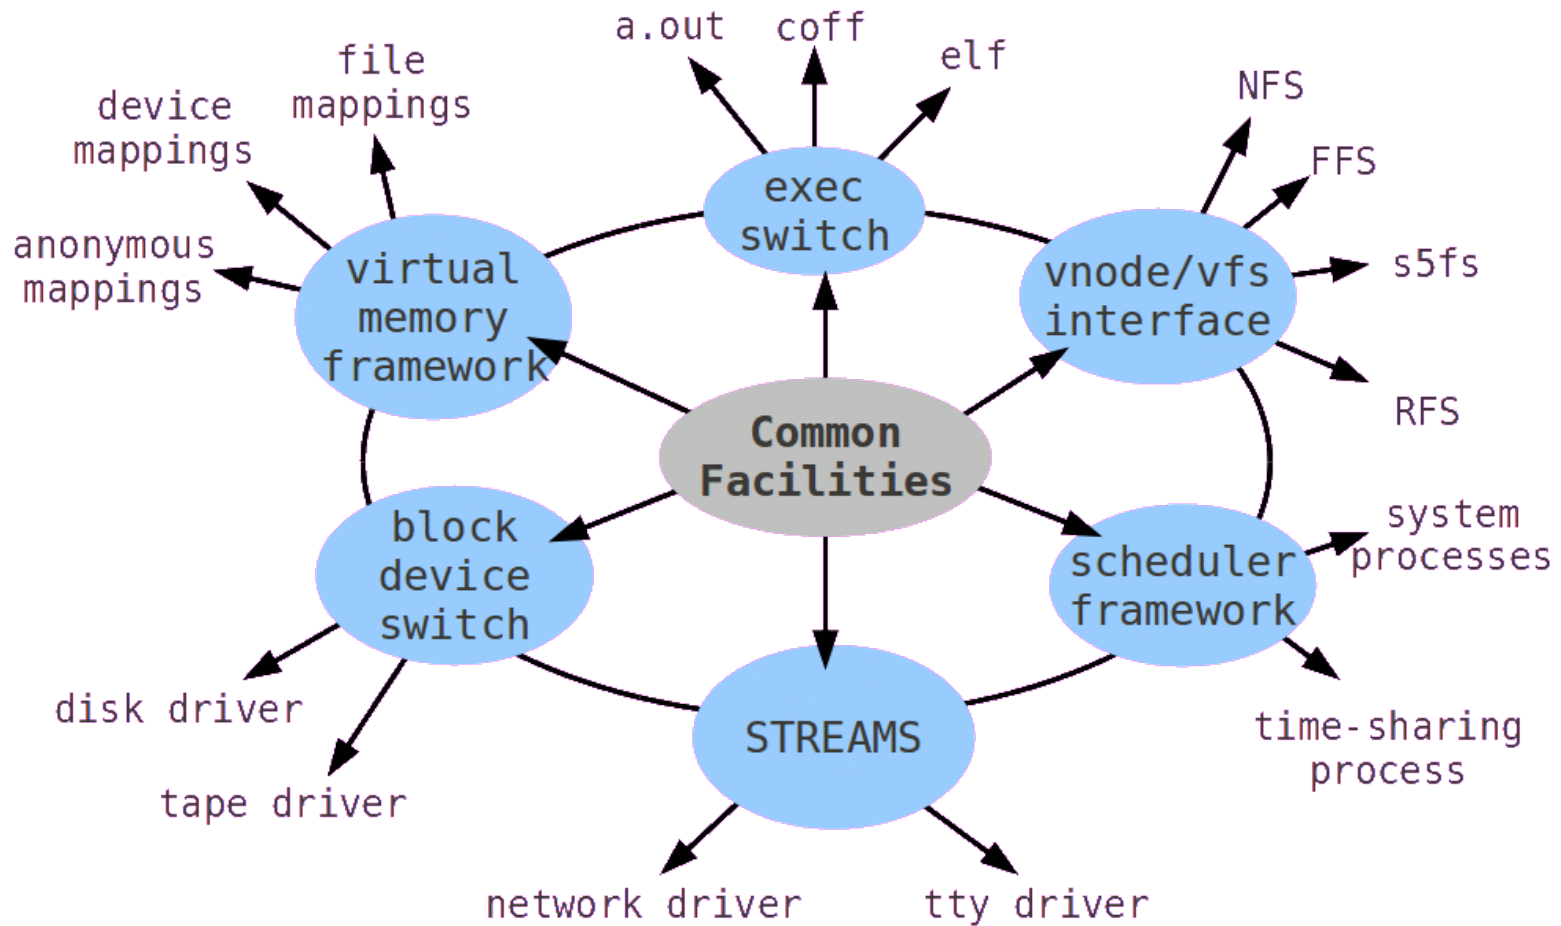
- Pro:
 - less context switches
 - no expensive communication
- Contra:
 - complications when exchanging functionality

- Microkernel

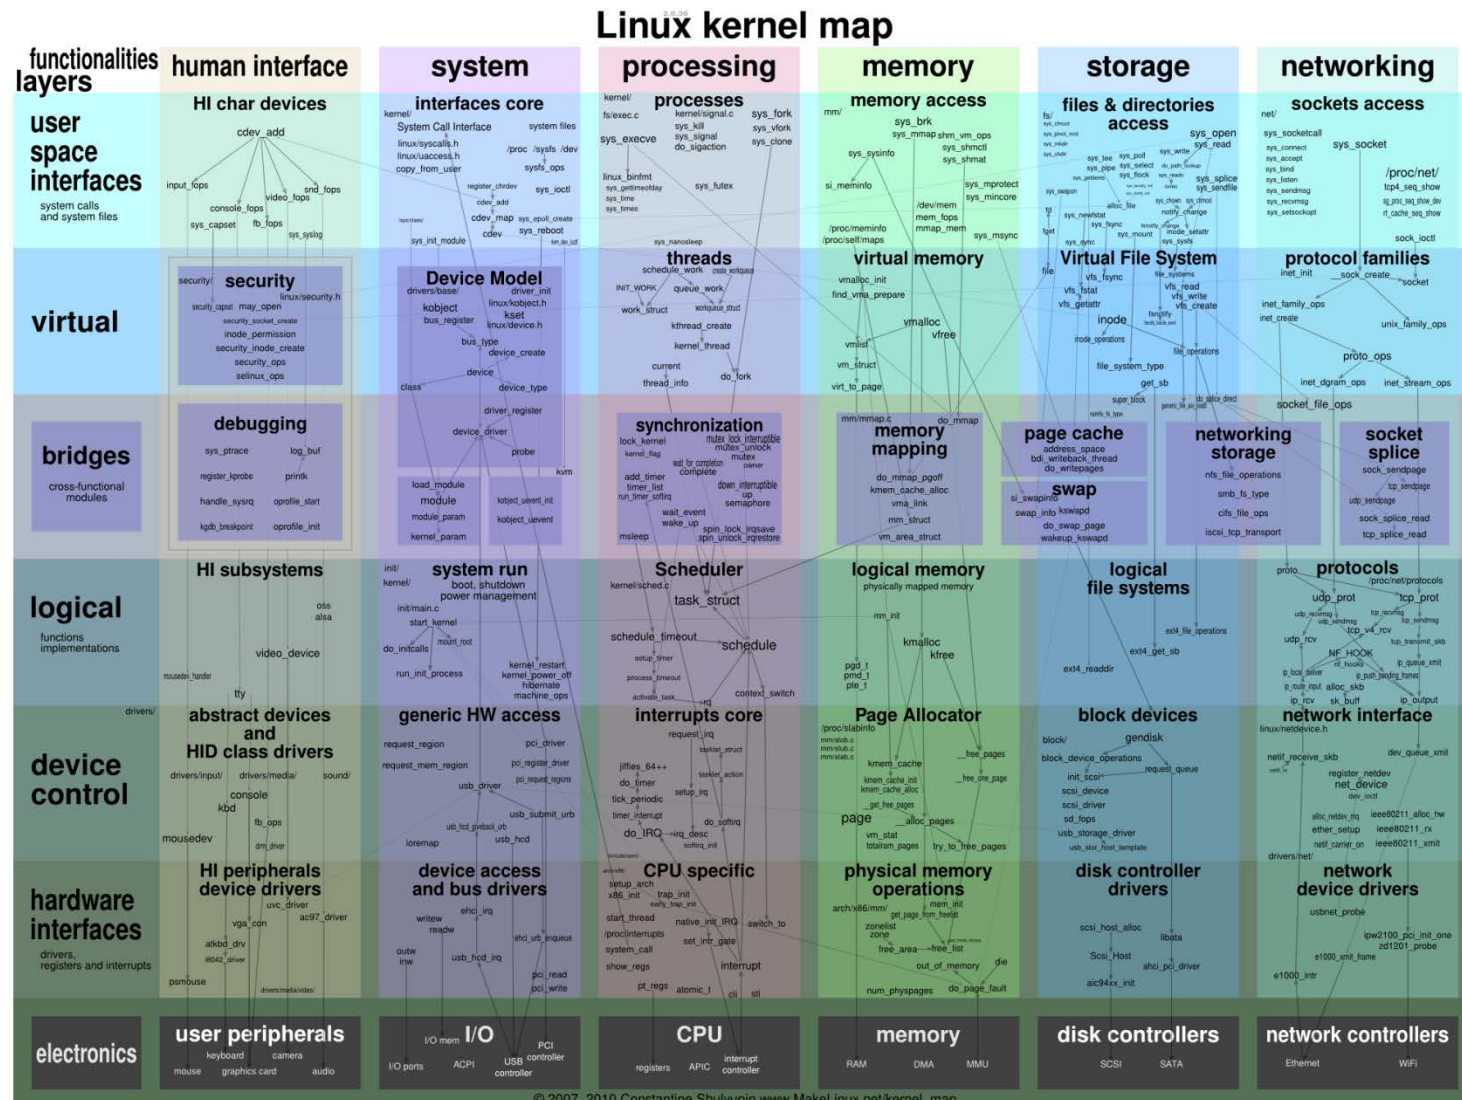


- Pro:
 - strict interfaces
 - less complexity, clear structure
- Contra:
 - speed
 - synchronization





Examples – Linux



Source: http://www.makelinux.net/kernel_map

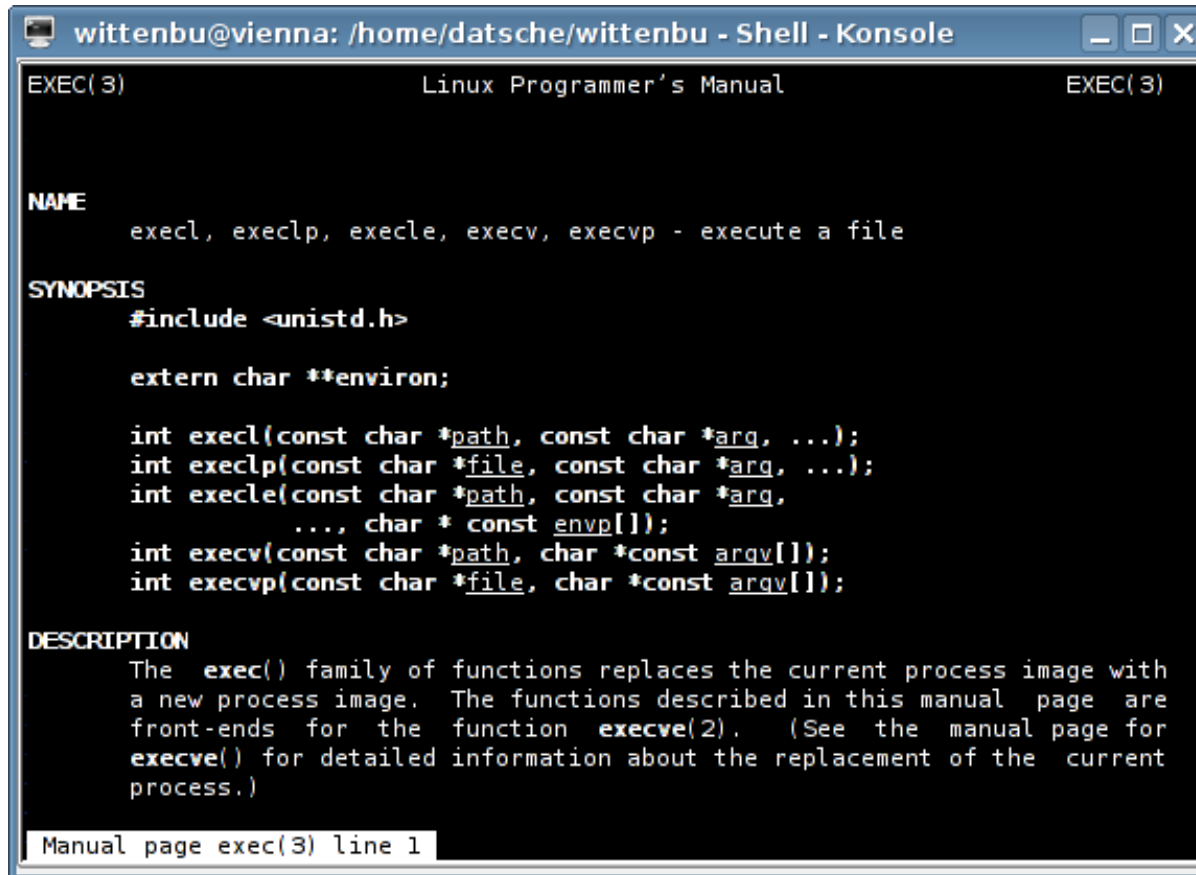
- Concurrency handling is outside the scope of this lecture.
 - Course “Nichtsequentielle Programmierung” in summer term

Some pointers/methods/ideas:

- In hardware:
 - Atomic operations:
 - ISA instructions that are guaranteed by design to run to completion
 - Interrupts:
 - Enable/disable interrupts via special ISA instructions
 - Allows other interrupt handlers to run to completion
- In software:
 - Spinlocks (busy waiting):
 - Short-term synchronization mechanism
 - Low overhead, avoid re-scheduling, wasteful on resources
 - Semaphores (wait queues):
 - Long-term synchronization mechanism
 - Synchronize for events on special purpose data structures

- UNIX-utility **man**

e.g. **man exec**



```
wittenbu@vienna: /home/datsche/wittenbu - Shell - Konsole
EXEC(3)                               Linux Programmer's Manual                               EXEC(3)

NAME
    execl, execlp, execl, execv, execvp - execute a file

SYNOPSIS
    #include <unistd.h>


    extern char **environ;

    int execl(const char *path, const char *arg, ...);
    int execlp(const char *file, const char *arg, ...);
    int execl(const char *path, const char *arg,
        ..., char * const envp[]);
    int execv(const char *path, char *const argv[]);
    int execvp(const char *file, char *const argv[]);

DESCRIPTION
    The exec() family of functions replaces the current process image with
    a new process image. The functions described in this manual page are
    front-ends for the function execve(2). (See the manual page for
execve() for detailed information about the replacement of the current
    process.)

Manual page exec(3) line 1
```

- Toolchain: set of programming tools that are used to build a product (executable)



- UNIX utility **top**:

```
wittenbu@vienna: /home/datsche/wittenbu - Shell - Konsole
top - 10:33:30 up 2 days, 1:04, 1 user, load average: 0.41, 0.26, 0.17
Tasks: 93 total, 1 running, 92 sleeping, 0 stopped, 0 zombie
Cpu(s):  0.3% us,  0.0% sy,  0.0% ni, 99.7% id,  0.0% wa,  0.0% hi,  0.0% si
Mem: 1033264k total,  967528k used,  65736k free, 160112k buffers
Swap: 2015992k total,    0k used, 2015992k free, 445496k cached

  PID USER      PR  NI  VIRT  RES  SHR  S  %CPU  %MEM    TIME+  COMMAND
11843 wittenbu  16   0  1944  968  740  R   0.3   0.1   0:00.54 top
   1 root      15   0  1584  520  452  S   0.0   0.1   0:01.44 init
   2 root       RT   0    0    0    0  S   0.0   0.0   0:00.00 migration/0
   3 root      34  19    0    0    0  S   0.0   0.0   0:00.00 ksoftirqd/0
   4 root      10  -5    0    0    0  S   0.0   0.0   0:02.72 events/0
   5 root      13  -5    0    0    0  S   0.0   0.0   0:00.02 khelper
   6 root      10  -5    0    0    0  S   0.0   0.0   0:00.00 kthread
   8 root      10  -5    0    0    0  S   0.0   0.0   0:00.15 kblockd/0
  11 root      10  -5    0    0    0  S   0.0   0.0   0:00.00 khubd
  13 root      10  -5    0    0    0  S   0.0   0.0   0:00.00 kseriod
 104 root      20   0    0    0    0  S   0.0   0.0   0:00.00 pdflush
 105 root      15   0    0    0    0  S   0.0   0.0   0:01.00 pdflush
 106 root      15   0    0    0    0  S   0.0   0.0   0:00.02 kswapd0
 107 root      20  -5    0    0    0  S   0.0   0.0   0:00.00 aio/0
 108 root      20  -5    0    0    0  S   0.0   0.0   0:00.00 xfslogd/0
 109 root      20  -5    0    0    0  S   0.0   0.0   0:00.00 xfsdatad/0
 764 root       11  -5    0    0    0  S   0.0   0.0   0:00.00 ata/0
 781 root       11  -5    0    0    0  S   0.0   0.0   0:00.00 kpsmoused
```


- Directories **/proc** and **/sys**
virtual directories that reflect general kernel behaviors
("everything is a file")

```
wittenbu@vienna: / - Shell - Konsole
wittenbu@vienna:/$ ls /proc/sys/kernel/
acct          modprobe      panic_on_oops  shmall
bootloader_type msgmax        pid_max        shmmax
cad_pid       msgmnb        printk         shmni
cap-bound     msgmni        printk_ratelimit sysrq
core_pattern  ngroups_max  printk_ratelimit_burst tainted
core_uses_pid osrelease     pty            threads-max
ctrl-alt-del  ostype       random         unknown_nmi_panic
domainname    overflowgid  randomize_va_space version
hostname      overflowuid  sem
hotplug       panic        sg-big-buff
wittenbu@vienna:/$ ls /proc/sys/vm
block_dump          legacy_va_layout    page-cluster
dirty_background_ratio lowmem_reserve_ratio percpu_pagelist_fraction
dirty_expire_centisecs max_map_count        swap_token_timeout
dirty_ratio          min_free_kbytes      swappiness
dirty_writeback_centisecs nr_pdflush_threads  vfs_cache_pressure
drop_caches          overcommit_memory
laptop_mode          overcommit_ratio
wittenbu@vienna:/$ ls /proc/sys/fs
aio-max-nr  dir-notify-enable  inode-state  mqueue  suid_dumpable
aio-nr      file-max           inotify      nfs      xfs
binfmt_misc file-nr            lease-break-time overflowgid
dentry-state inode-nr           leases-enable overflowuid
wittenbu@vienna:/$
```

1. Introduction and Motivation

2. Subsystems, Interrupts and System Calls

3. Processes

4. Memory

5. Scheduling

6. I/O and File System

7. Booting, Services, and Security