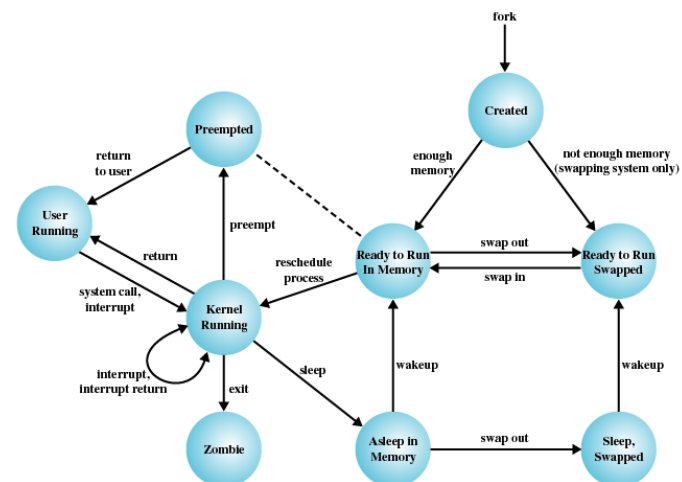




# Operating Systems & Computer Networks

## Processes



1. Introduction and Motivation
2. Subsystems, Interrupts and System Calls
- 3. Processes**
4. Memory
5. Scheduling
6. I/O and File System
7. Booting, Services, and Security

# Definitions of a Process


- Program in execution
- Instance of a program running on a computer
  - There may be multiple instances of the same program, each as a separate process
- Unit characterized by
  - Execution of a sequence of instructions
  - Current state
  - Associated block of memory

# Related Concepts to “Process”

- Thread: One (of several) runtime entities that share the same address space
  - Easy cooperation, requires explicit synchronization
  - A process may consist of several threads
- Application: User-visible entity, one or more processes

# Program vs. Process

- Multiple parts
  - Program code → text section
  - Current activity → program counter, processor registers
  - Stack → temporary data
  - Data section → global variables
  - Heap → dynamic memory
- Program is passive entity, process is active
  - Program becomes process when executable file loaded into memory
- One program can be several processes



- Interleaved execution (by scheduling) of multiple processes
  - Maximization of processor utilization
  - Reduction of response time
- Allocation of resources for processes
  - Consideration of priorities
  - Avoidance of deadlocks
- Support for Inter-Process Communication (IPC)
- On-demand user-level process creation
  - Structuring of applications

# Process execution (Trace)

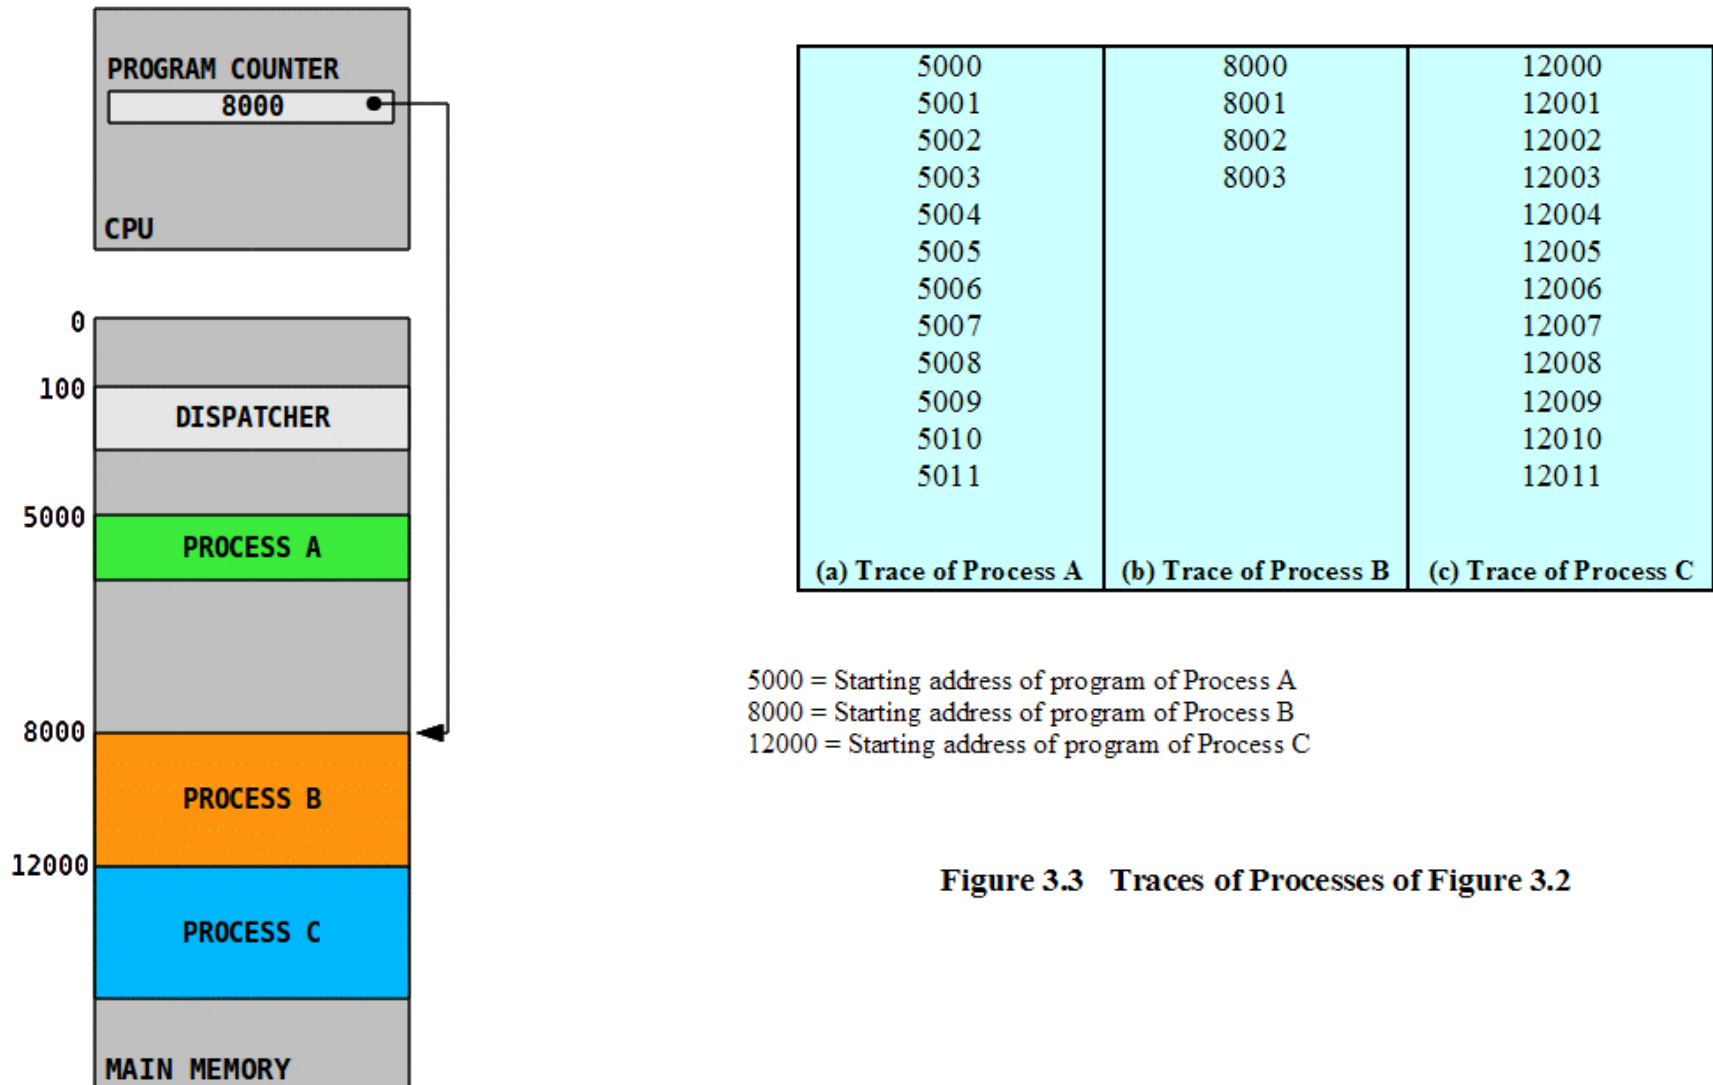
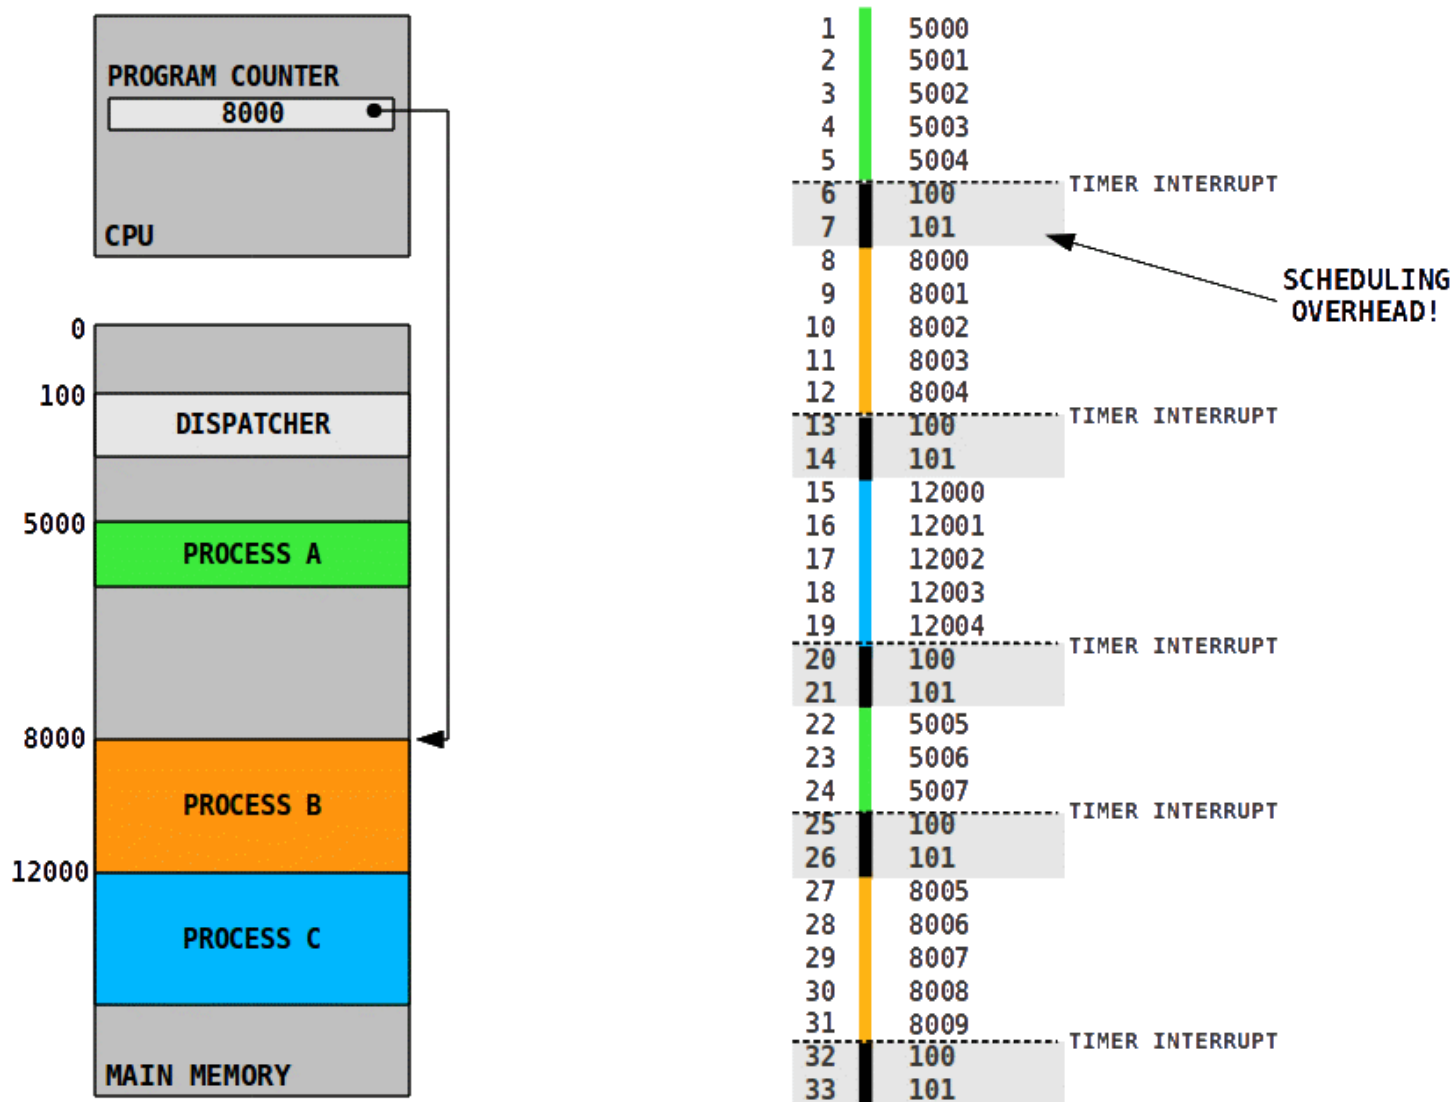


Figure 3.3 Traces of Processes of Figure 3.2




# Process execution (Trace)





# Simple Process Model

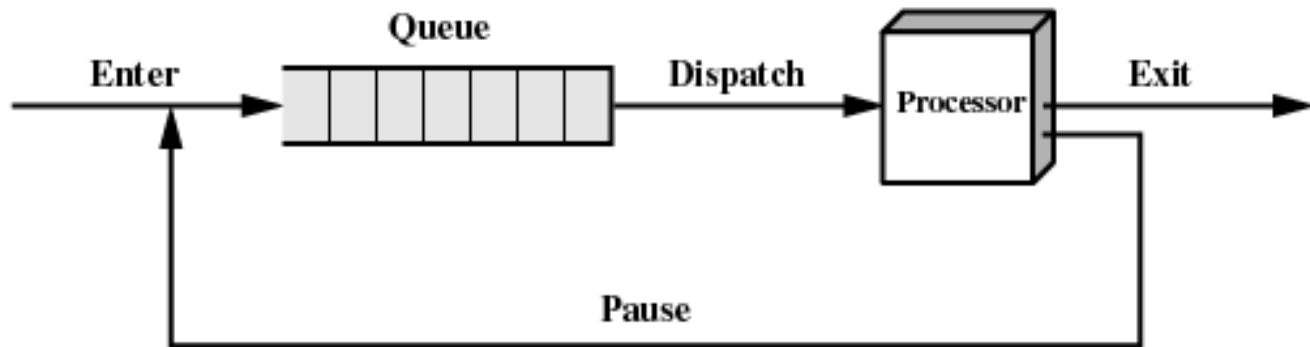
- Process is in one of two states:
  - running
  - not running



- How to implement?

# Simple Process Model

- Running processes managed in queue:



- What information required?

# Process Control Block (PCB)

**Definition:** OS data structure which contains the information needed to manage a process (one PCB per process)

|                     |  |
|---------------------|--|
| Process identifiers | <ul style="list-style-type: none"><li>• IDs of process, parent process, and user</li></ul>   |
| CPU state           | <ul style="list-style-type: none"><li>• User-visible registers</li><li>• Control and status registers:<ul style="list-style-type: none"><li>• Stack pointer (SP)</li><li>• Program counter (PC)</li><li>• Processor status word (PSW)</li></ul></li></ul>  |
| Control information | <ul style="list-style-type: none"><li>• Scheduling information:<ul style="list-style-type: none"><li>• Process state, priority, awaited event</li></ul></li><li>• Accounting information:<ul style="list-style-type: none"><li>• Amount of memory used, CPU time elapsed</li></ul></li><li>• Memory management:<ul style="list-style-type: none"><li>• Location and access state of all user data</li></ul></li><li>• I/O management:<ul style="list-style-type: none"><li>• Devices currently opened (files, sockets)</li></ul></li></ul> |

# Process Control Block (PCB)

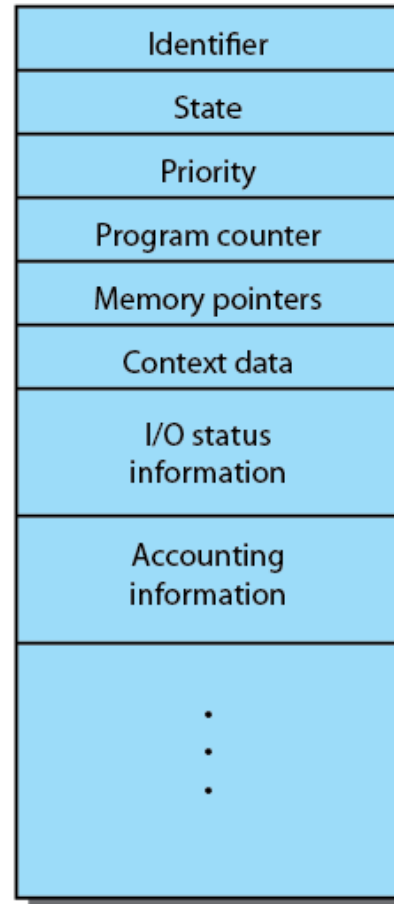


Figure 3.1 Simplified Process Control Block


# Reasons for Process Creation

- Interactive logon
  - User logs onto a terminal
    - May create several processes as part of logon procedure (e.g. GUI)
- Created by the OS to provide a service
  - Provide a service to user program in the background (e.g. printer spooling)
    - Either at boot time or dynamically in response to requests (e.g. HTTP)
- Spawned at application start-up
  - Separation of a program into separate processes for algorithmic purposes
- Always spawned by existing process
  - Operating system creates first process at boot time
  - Processes are organized in a tree-like structure (``ps tree``)

# Process Termination

- Execution of process is completed
    - process terminates itself by system call
  - Other user process terminates the process
    - Parent process or other authorized processes
  - OS terminates process for protection reasons
    - Invalid instruction (process tries to execute data)
    - Privileged instruction in user mode
    - Process tries to access memory without permission
    - I/O-Error
    - Arithmetic error
- Some exceptions can be caught and handled by the process.

- Simple model with two states

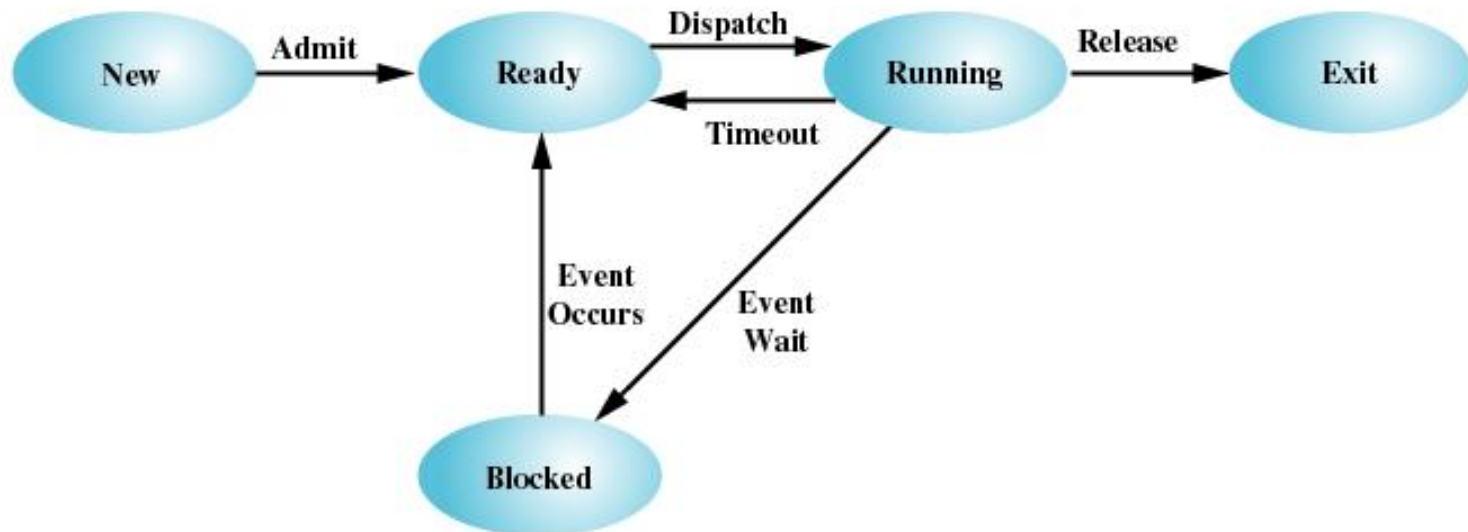


- Problems
  - Most of the processes will be waiting for IO
  - Different IO devices
  - Different priorities

➤ Extend the model


Five states including creation, termination, and resource handling:

- **Running:** currently being executed
- **Ready:** ready to run, waiting for execution
- **Blocked:** not ready to run, waiting for external event, e.g., completion of I/O operation
- **New:** newly created process, not yet in running set
- **Exit:** completed/terminated process, removed from running set



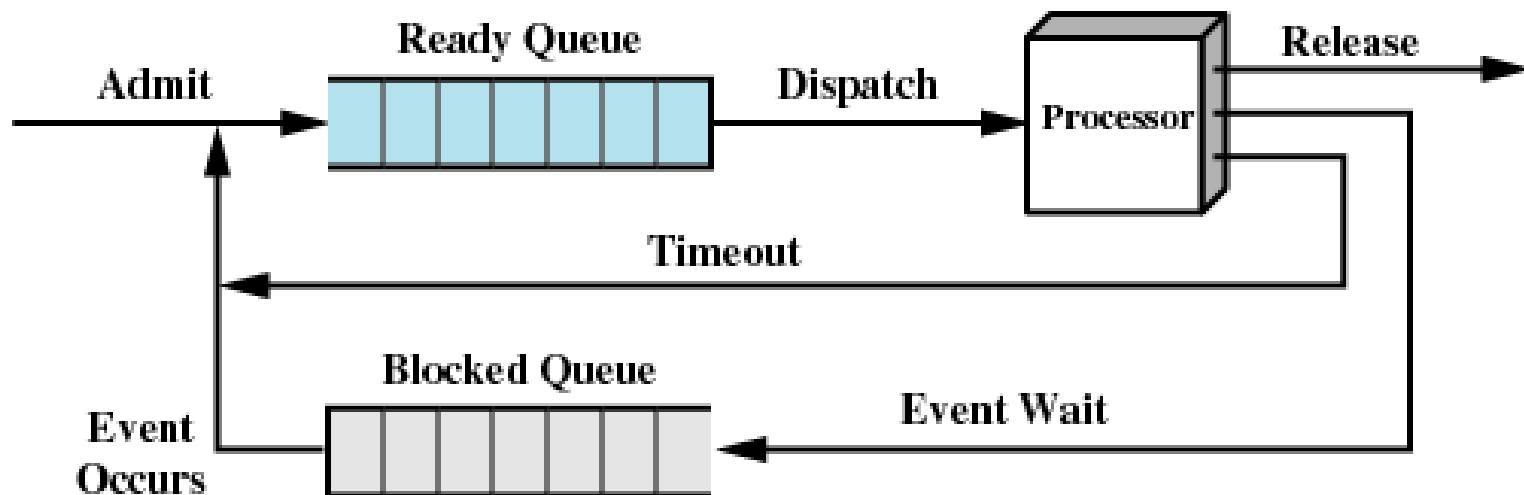


# Process States over Time



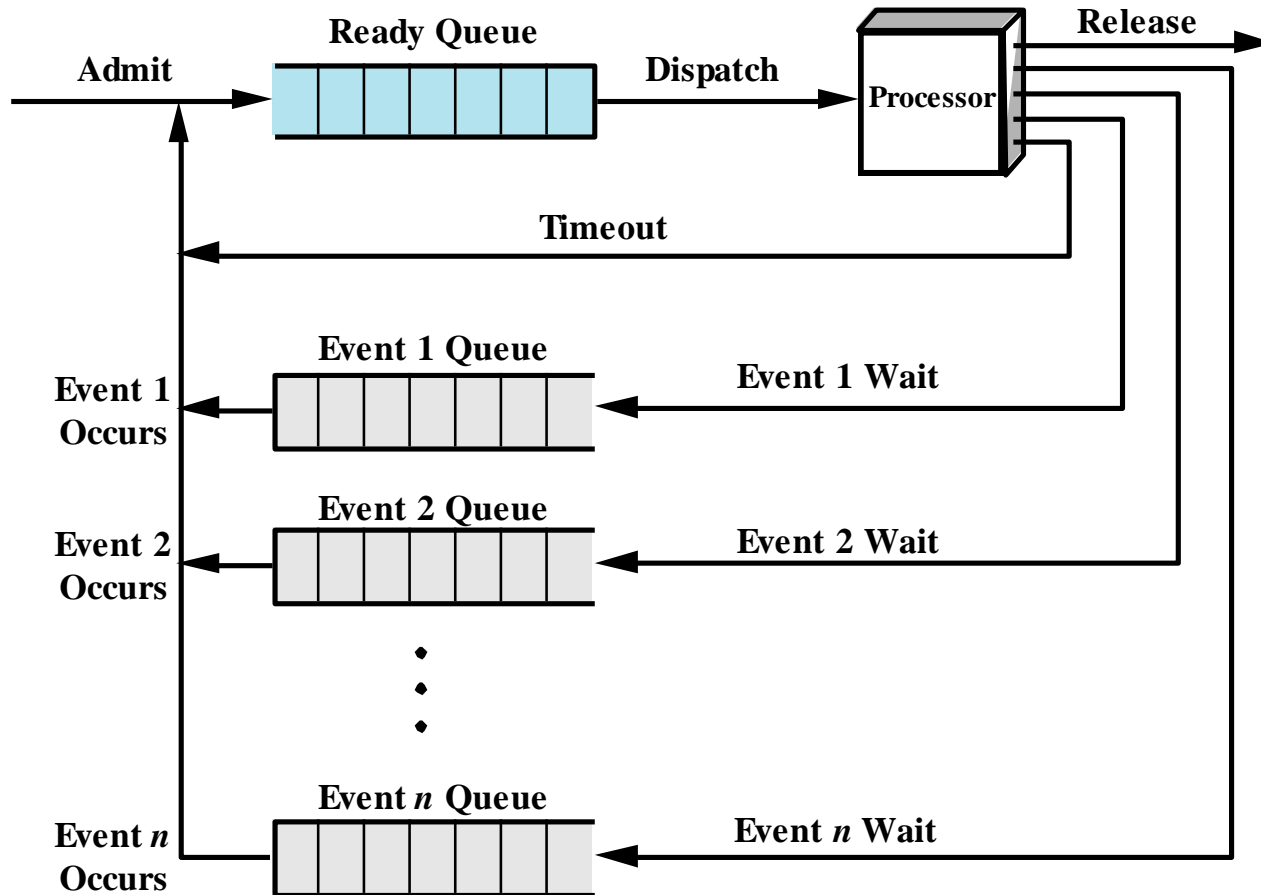
# Implementation of Process States

- Assign process to different queues based on state of required resources
- Two queues:
  - Ready processes (all resources available)
  - Blocked processes (at least one resource busy)



➤ But what happens if processes need *different* resources?

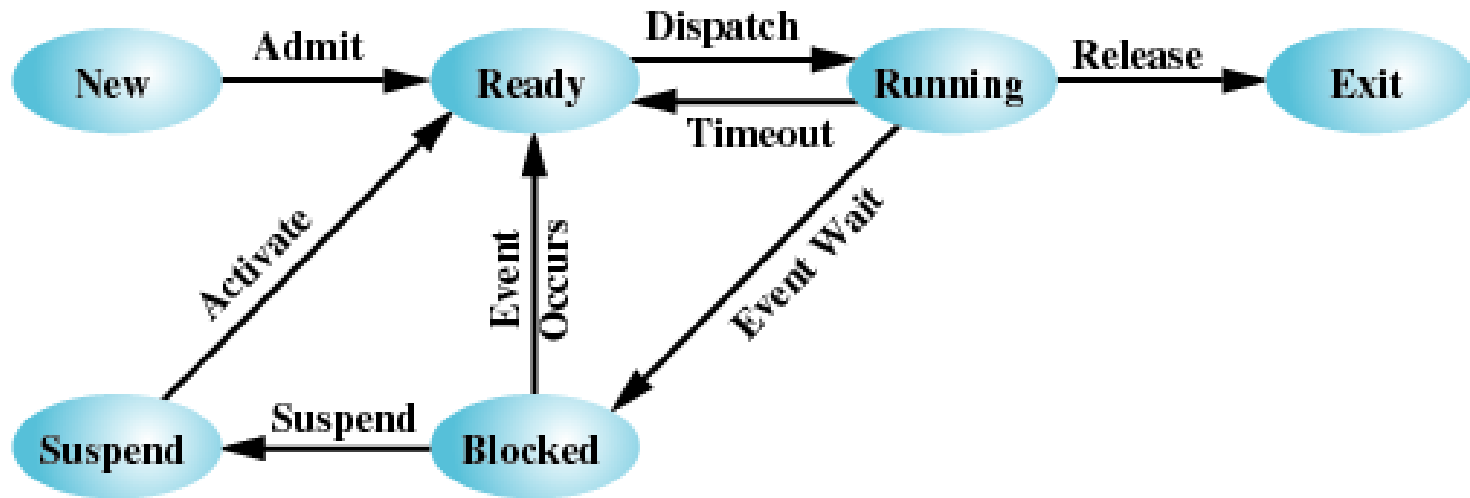
- Several queues one for each resource / type of resource



➤ More efficient, but fairness issues must be considered

Swapping motivated by two observations:

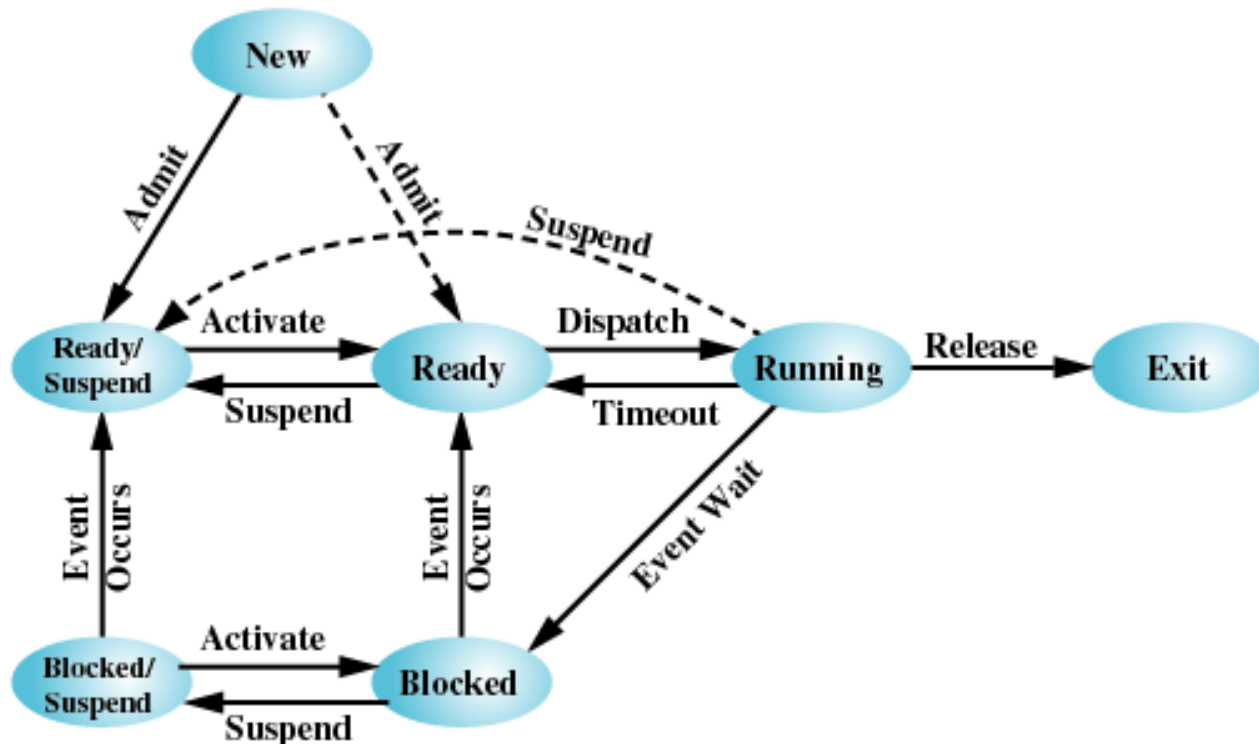
- Physical main memory is (was) a scarce resource
  - Blocked processes may wait for longer periods of time (e.g. during I/O, while waiting for requests, ...)
- Swap blocked processes to secondary storage thereby reducing memory usage



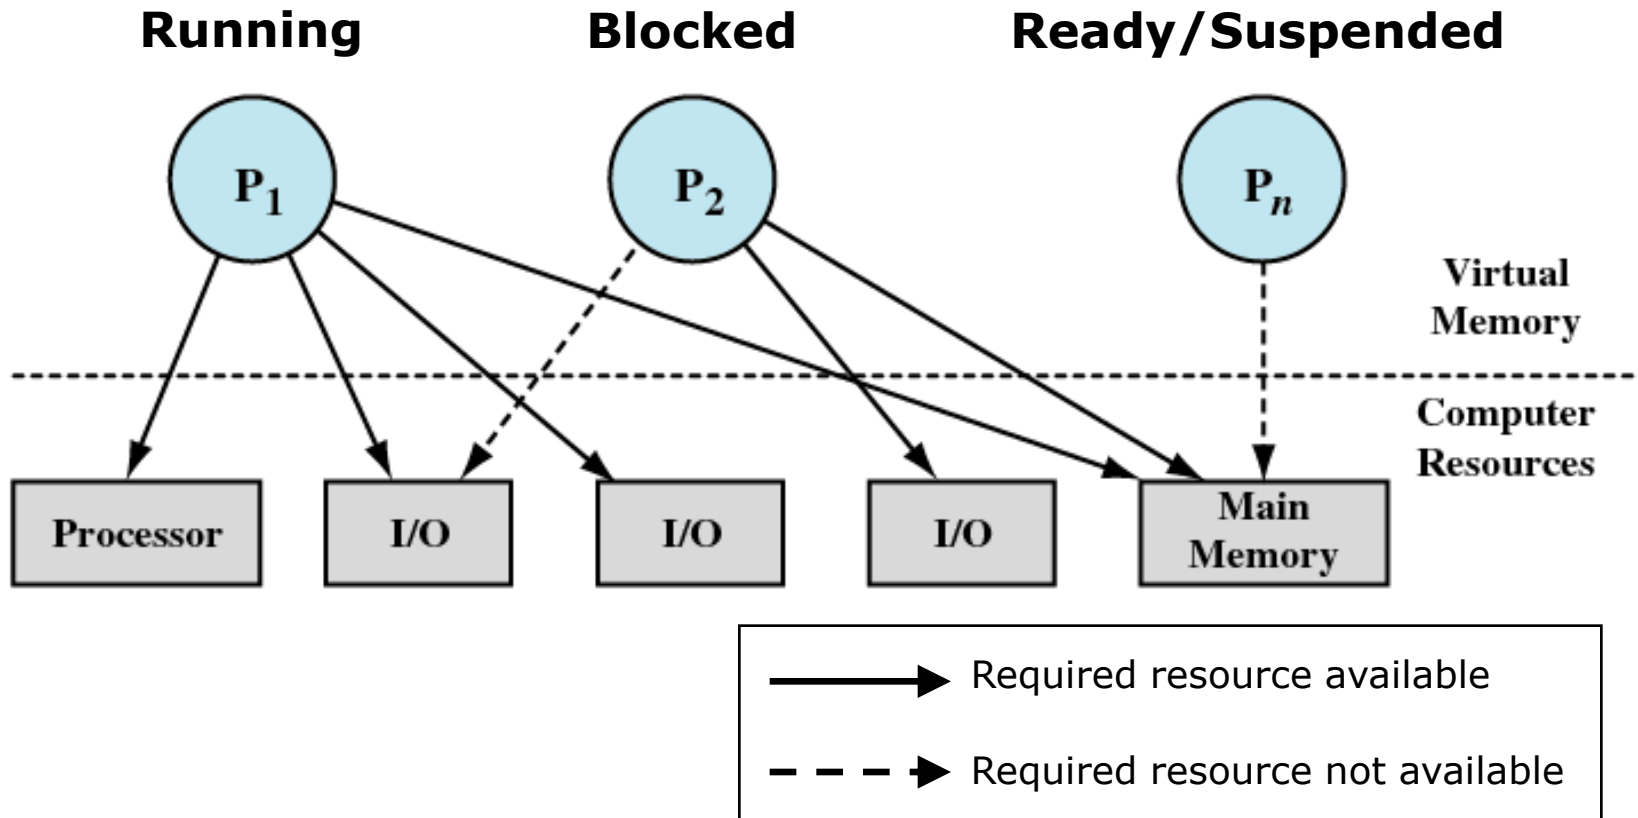
# Extended Process State Diagram

Two additional considerations:

- Blocked/swapped processes may become ready to run when event occurs
- Ready and/or running processes may be swapped even without waiting for event



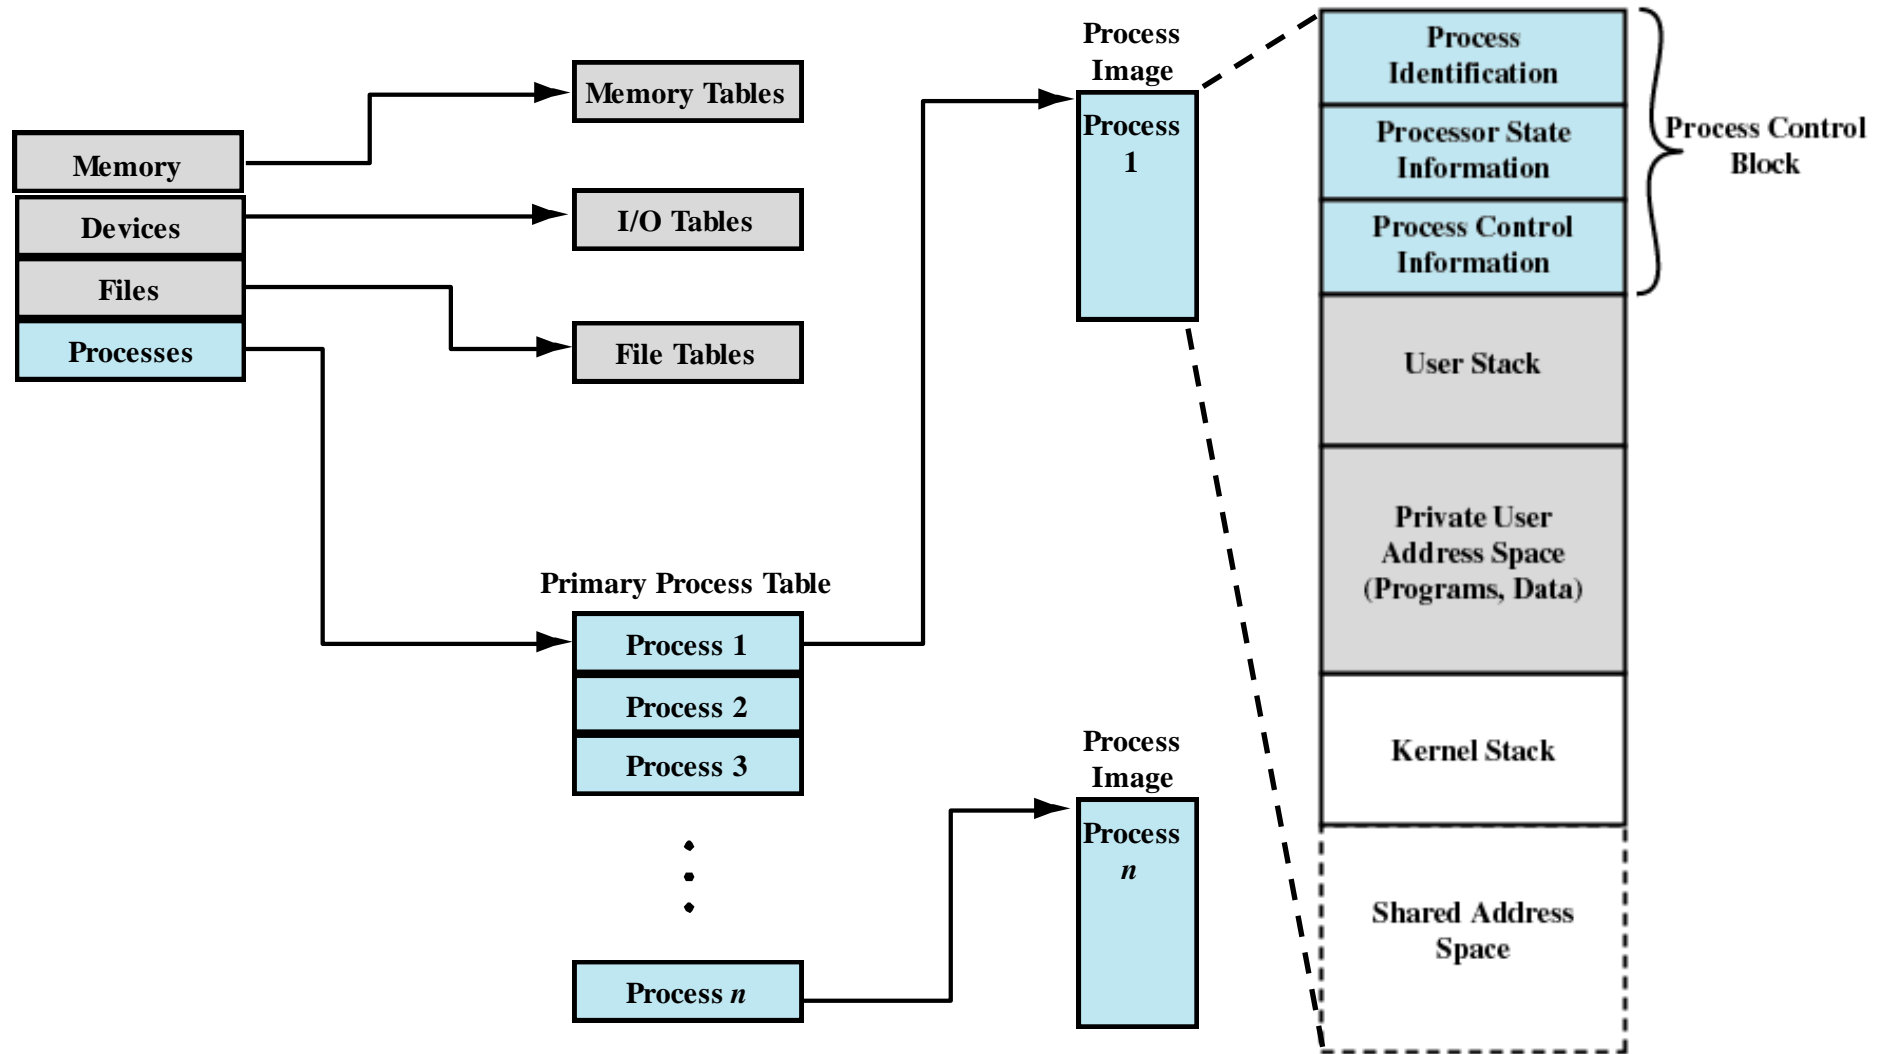
- Process state reflects allocated resources:



## Global data structures for processes and resources usage:


- Process tables:
  - Process Control Block (PCB)
  - Location of process image in memory
  - Resources (process-specific view)
- Memory tables:
  - Allocation of primary and secondary memory
  - Protection attributes of blocks of (shared) memory
  - Virtual memory management
- I/O tables:
  - Allocation of I/O devices, assignment to processes
  - State of current operation and corresponding memory region
- File tables:
  - Currently open files
  - Location on storage media / secondary memory
  - State and attributes

# Process Control Table and Image






- Separated kernel and processes:
  - Separate memory and stack for kernel
  - Kernel is no process
  - Expensive and unsafe




(a) Separate kernel

- Execution of system calls as part of user process, but in kernel mode:
  - Kernel functions use same address space
  - Same process switches into privileged mode (Ring 0)
  - Less expensive and quite safe



**(b) OS functions execute within user processes**


- Microkernel:
  - Collection of system processes that provide OS services
  - Quite expensive but very safe



**(c) OS functions execute as separate processes**

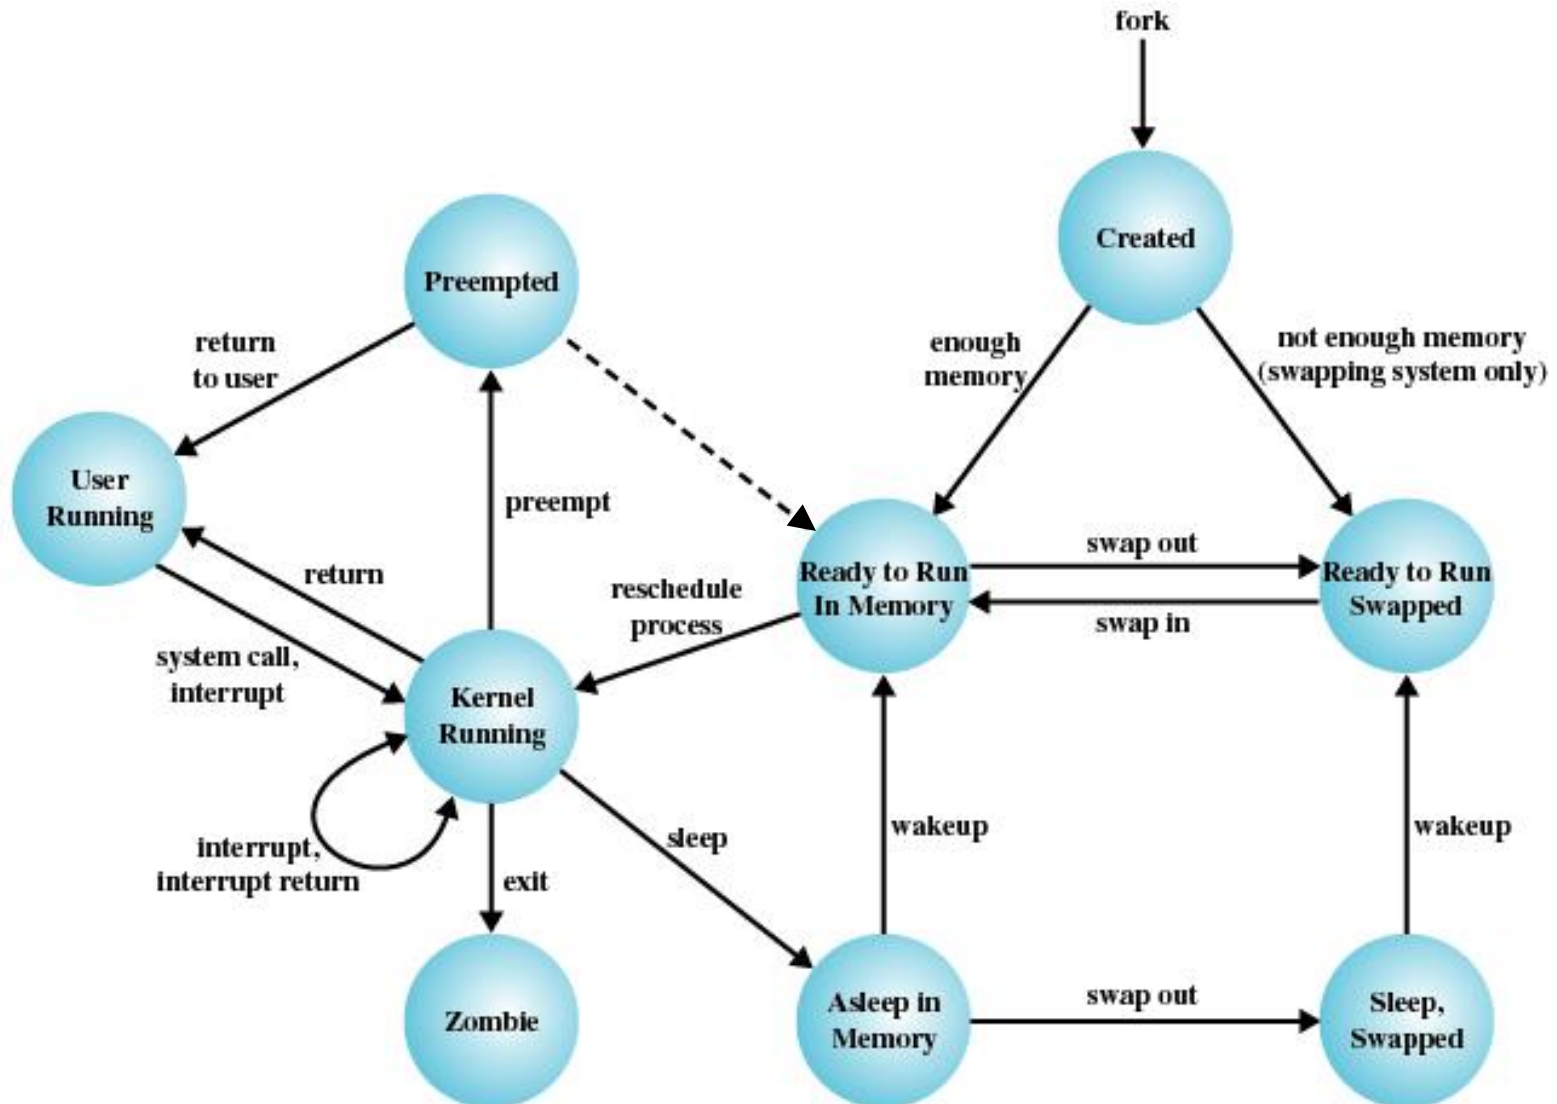
# Example: UNIX – Architecture

- Process architecture that executes kernel functions in the context of a user process



- Two modes are used: user / kernel mode (Ring 3/Ring 0)
- Two types of processes: system / user processes
- System processes are implemented as part of kernel to run background services, e.g. swapping

# Example: UNIX – Process State Diagram



# Example: UNIX – Process States

|                                |  |
|--------------------------------|--|
| <b>User Running</b>            | Executing in user mode.  |
| <b>Kernel Running</b>          | Executing in kernel mode.  |
| <b>Ready to Run, in Memory</b> | Ready to run as soon as the kernel schedules it.   |
| <b>Asleep in Memory</b>        | Unable to execute until an event occurs; process is in main memory (a blocked state).  |
| <b>Ready to Run, Swapped</b>   | Process is ready to run, but the swapper must swap the process into main memory before the kernel can schedule it to execute.    |
| <b>Sleeping, Swapped</b>       | The process is awaiting an event and has been swapped to secondary storage (a blocked state).                                    |
| <b>Preempted</b>               | Process is returning from kernel to user mode, but the kernel preempts it and does a process switch to schedule another process. |
| <b>Created</b>                 | Process is newly created and not yet ready to run.   |
| <b>Zombie</b>                  | Process no longer exists, but it leaves a record for its parent process to collect.  |

# Related System Calls

- **`int exeve(const char *filename, char *const argv[], char *const envp[])`**
  - Executes program pointed to by **`filename`** with arguments **`argv`** and environment **`envp`** (in the form of key=value)
  - Effectively replaces the current program with another one
  - **`exec()`** family of library function
- **`pid_t fork(void)`**
  - Creates child process that differs from parent only in its PID (process identifier) and PPID (parent process identifier)
  - Returns 0 for child process and child's PID for parent process
- **`void _exit(int status)`**
  - Terminates calling process; closes open file descriptors; children are adopted by process 1; signals termination to parent
  - **`exit()`** library function
- **`pid_t wait(int *status)`**
  - Wait for state change in child of calling process

# Programming Example

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>

main()
{
    int status;
    pid_t pid;

    pid = fork();
    if(pid == 0) {
        printf("Child process running...\n");
        // Do something...
        printf("Child process done.\n");
        exit(123);
    }
    else if(pid > 0) {
        printf("Parent process, waiting for child %d...\n", pid);
        pid = wait(&status);
        printf("Child process %d terminated, status %d.\n", pid, WEXITSTATUS(status));
        exit(EXIT_SUCCESS);
    }
    else {
        printf("fork() failed\n");
        exit(EXIT_FAILURE);
    }
}
```



```
wittenbu@vienna: /home/datsche/wittenbu - Shell - Konsole
wittenbu@vienna:~$ ps
  PID TTY          TIME CMD
 19047 pts/1        00:00:00 csh
  19050 pts/1        00:00:00 bash
  19243 pts/1        00:00:00 ps
wittenbu@vienna:~$ kword &
[1] 19244
wittenbu@vienna:~$ ps
  PID TTY          TIME CMD
 19047 pts/1        00:00:00 csh
  19050 pts/1        00:00:00 bash
  19244 pts/1        00:00:01 kword
  19245 pts/1        00:00:00 ps
wittenbu@vienna:~$ kill -9 19244
wittenbu@vienna:~$ ps
  PID TTY          TIME CMD
 19047 pts/1        00:00:00 csh
  19050 pts/1        00:00:00 bash
  19246 pts/1        00:00:00 ps
[1]+  Killed                  kword
wittenbu@vienna:~$
```

# User-Level Process Control

Activity Monitor (All Processes)

CPU Memory Energy Disk Network

| Process Name         | % CPU | CPU Time   | Threads | Idle Wake Ups | PID   | User    |
|----------------------|-------|------------|---------|---------------|-------|---------|
| WindowServer         | 14,8  | 41:09,74   | 6       | 23            | 180   | _window |
| Activity Monitor     | 7,2   | 8,04       | 5       | 0             | 76525 | gunes   |
| sysmond              | 4,0   | 20,04      | 3       | 1             | 210   | root    |
| Dock                 | 2,6   | 16:59,24   | 12      | 4             | 621   | gunes   |
| VShieldScanManager   | 1,1   | 4:03:04,70 | 11      | 147           | 65    | root    |
| kernel_task          | 0,9   | 1:11:38,36 | 104     | 213           | 0     | root    |
| cma                  | 0,8   | 23:58,24   | 20      | 49            | 330   | root    |
| VShieldScanner       | 0,8   | 14:09,20   | 3       | 1             | 402   | root    |
| VShieldScanner       | 0,7   | 14:00,74   | 3       | 0             | 400   | root    |
| Microsoft PowerPoint | 0,7   | 1:47,16    | 11      | 4             | 75735 | gunes   |
| Mail                 | 0,6   | 3:46,47    | 28      | 0             | 68888 | gunes   |
| VShieldScanner       | 0,6   | 14:01,51   | 3       | 0             | 401   | root    |
| SystemUIServer       | 0,5   | 35,17      | 10      | 1             | 622   | gunes   |
| pri_disp_service     | 0,5   | 17:54,57   | 18      | 1             | 444   | root    |
| Google Chrome        | 0,5   | 10:50,47   | 41      | 19            | 94061 | gunes   |
| launchservicesd      | 0,4   | 22,55      | 9       | 0             | 51    | root    |
| coreservicesd        | 0,3   | 36,89      | 5       | 0             | 71    | root    |
| Dashboard            | 0,3   | 8:47,49    | 13      | 1             | 1675  | gunes   |
| Finder               | 0,3   | 2:26,25    | 18      | 0             | 623   | gunes   |

|         |         |          |            |     |
|---------|---------|----------|------------|-----|
| System: | 2,10 %  | CPU LOAD | Threads:   | 996 |
| User:   | 4,20 %  |          | Processes: | 197 |
| Idle:   | 93,70 % |          |            |     |

Microsoft PowerPoint (75735)

Parent Process: [launchd \(603\)](#) User: guenes (503)

Process Group: Microsoft PowerPoint

% CPU: 8,68 Recent hangs: 0

Memory Statistics Open Files and Ports

|                      |          |
|----------------------|----------|
| Real Memory Size:    | 178,4 MB |
| Virtual Memory Size: | 1,36 GB  |
| Shared Memory Size:  | 63,2 MB  |
| Private Memory Size: | 51,3 MB  |

Sample Quit

1. Introduction and Motivation
2. Subsystems, Interrupts and System Calls
- 3. Processes**
4. Memory
5. Scheduling
6. I/O and File System
7. Booting, Services, and Security