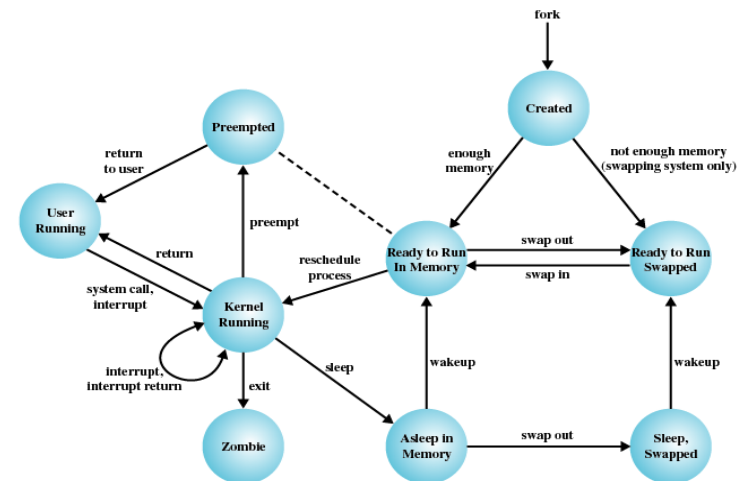


# Operating Systems & Computer Networks

## Scheduling

Types of Scheduling  
Decision Modes  
Process Priorities  
Scheduling Policies

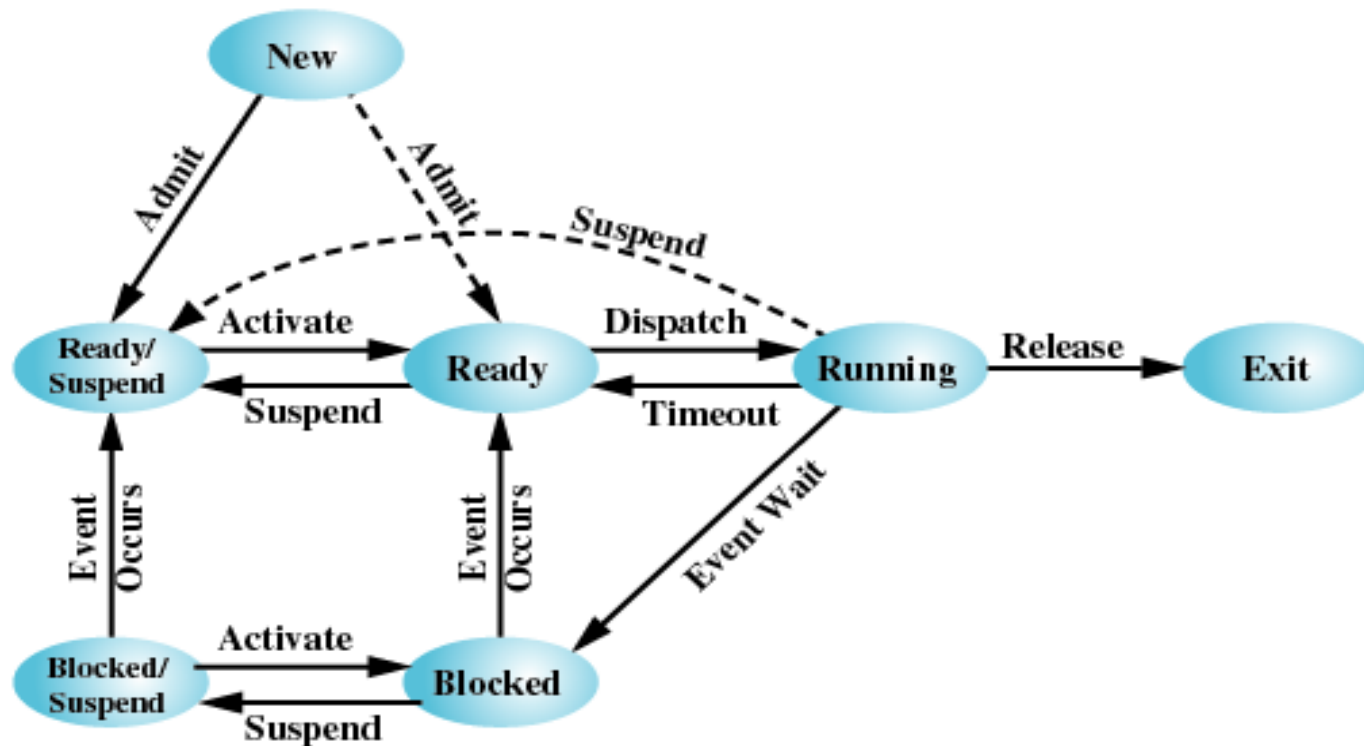


1. Introduction and Motivation
2. Subsystems, Interrupts and System Calls
3. Processes
4. Memory
- 5. Scheduling**
6. I/O and File System
7. Booting, Services, and Security

# Definition and Goals

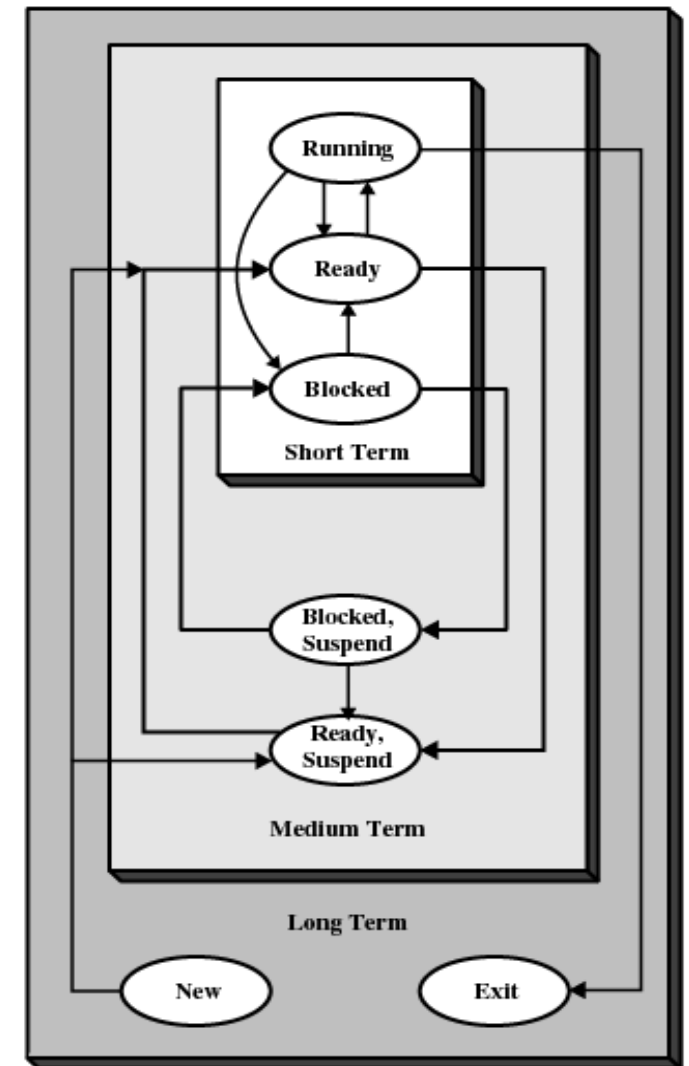
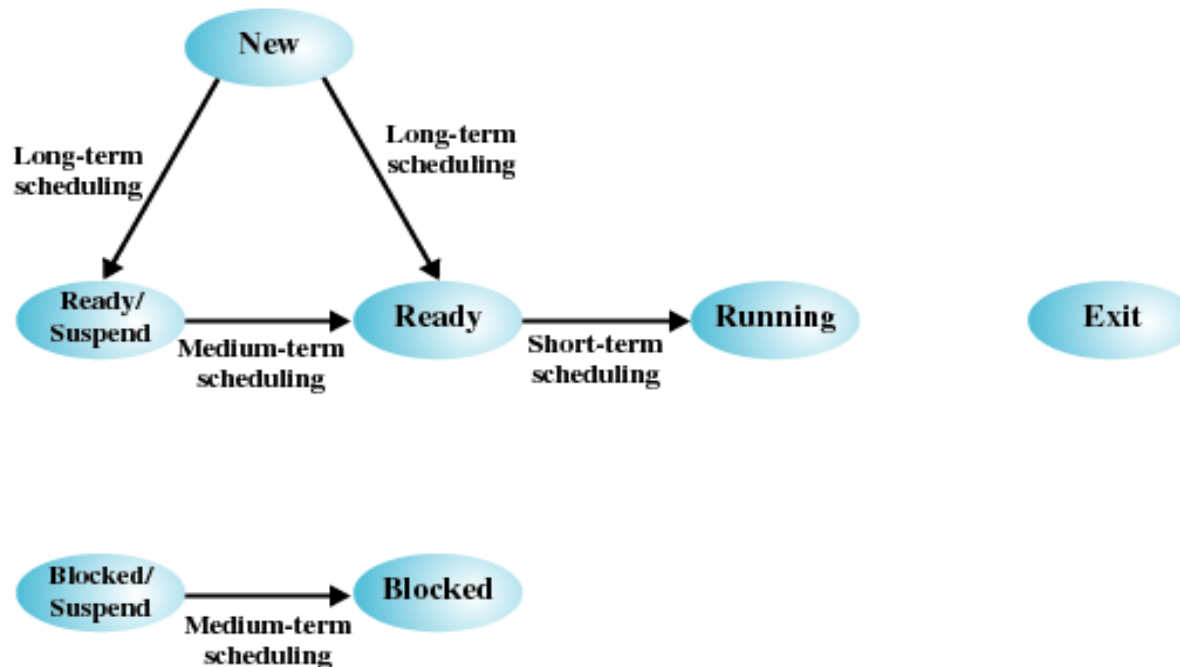
- Assign processes to be executed by the processor(s)
- More general: Assign consumers to resources
  - Examples: I/O requests → Device-specific queues  
Memory pages → Primary/secondary memory
- Goals:
  - Throughput, i.e., effectively use processing time
  - Response time / fairness, i.e., interactivity of individual processes
  - Processor efficiency, i.e., optimal utilization of CPU (as resource)
- Conflicting goals: Maximal throughput means unpredictable response time (and vice versa)

- Scheduling decisions correspond to state transitions in process state graph



# Process States and Scheduling

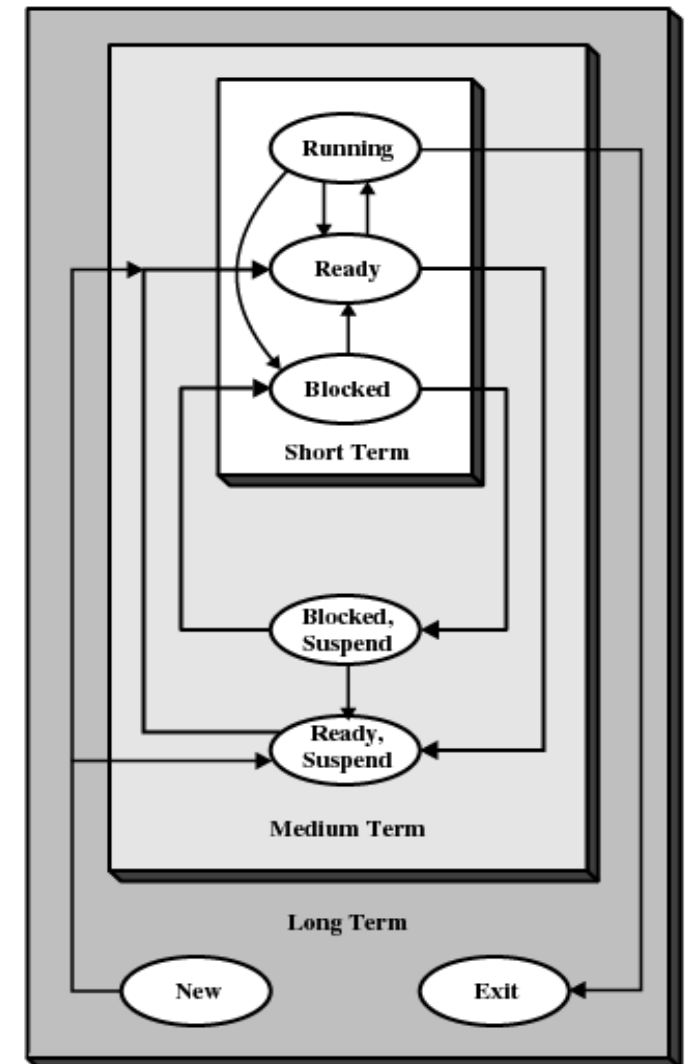
- Scheduling decisions correspond to state transitions in process state graph
- States form hierarchy depending on transition frequency
  - Import to consider when choosing and implementing scheduling algorithms



# Types of Scheduling

## Long-term scheduling

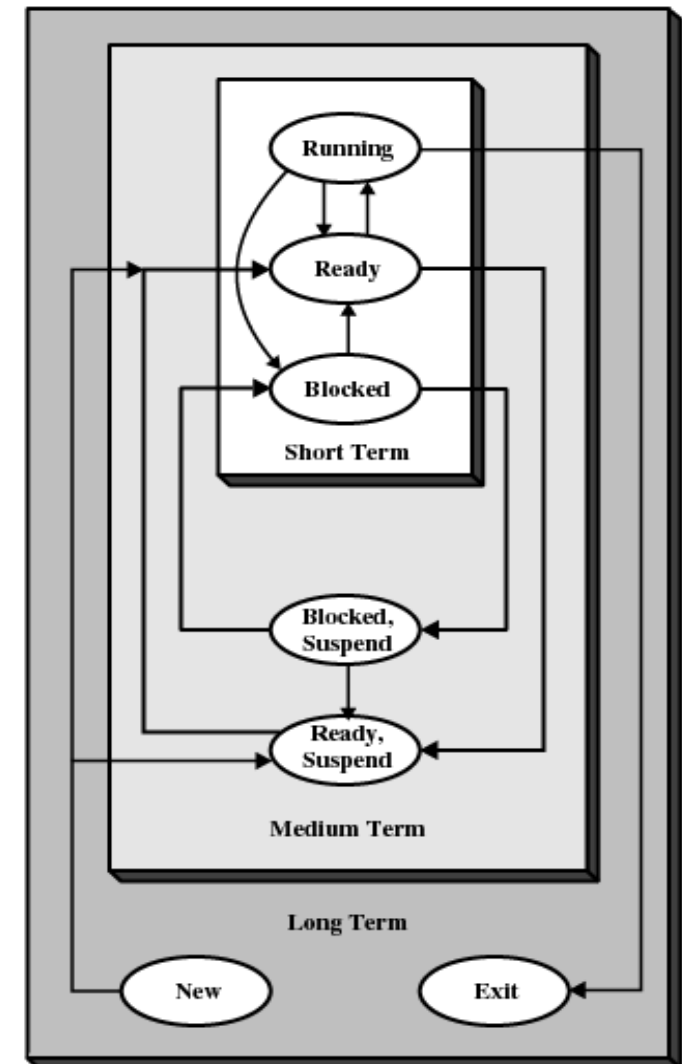
- Whether to add process to running queue and execute it
  - Determines which programs are admitted to system for processing, e.g., based on user
  - Specifies degree of multiprogramming, i.e., maximal number of processes
  - The more processes, the smaller percentage of time each process is executed
- How many processes should be allowed?



# Types of Scheduling

## Medium-term scheduling


- Whether to add/remove existing process (that is only partially in primary memory)
  - Part of swapping function
  - Based on need to dynamically manage degree of multiprogramming (considering available resources)
- Should processes be swapped in or out? If so, which ones?



# Types of Scheduling

## Short-term scheduling

- Which one of fully available processes to run
  - Known as “dispatcher”
  - Executes most frequently
    - Overhead / algorithmic complexity matters
  - Invoked when event occurs (clock interrupts, I/O interrupts, operating system calls, signals)
- Whose turn is it?






# Types of Scheduling

## I/O scheduling

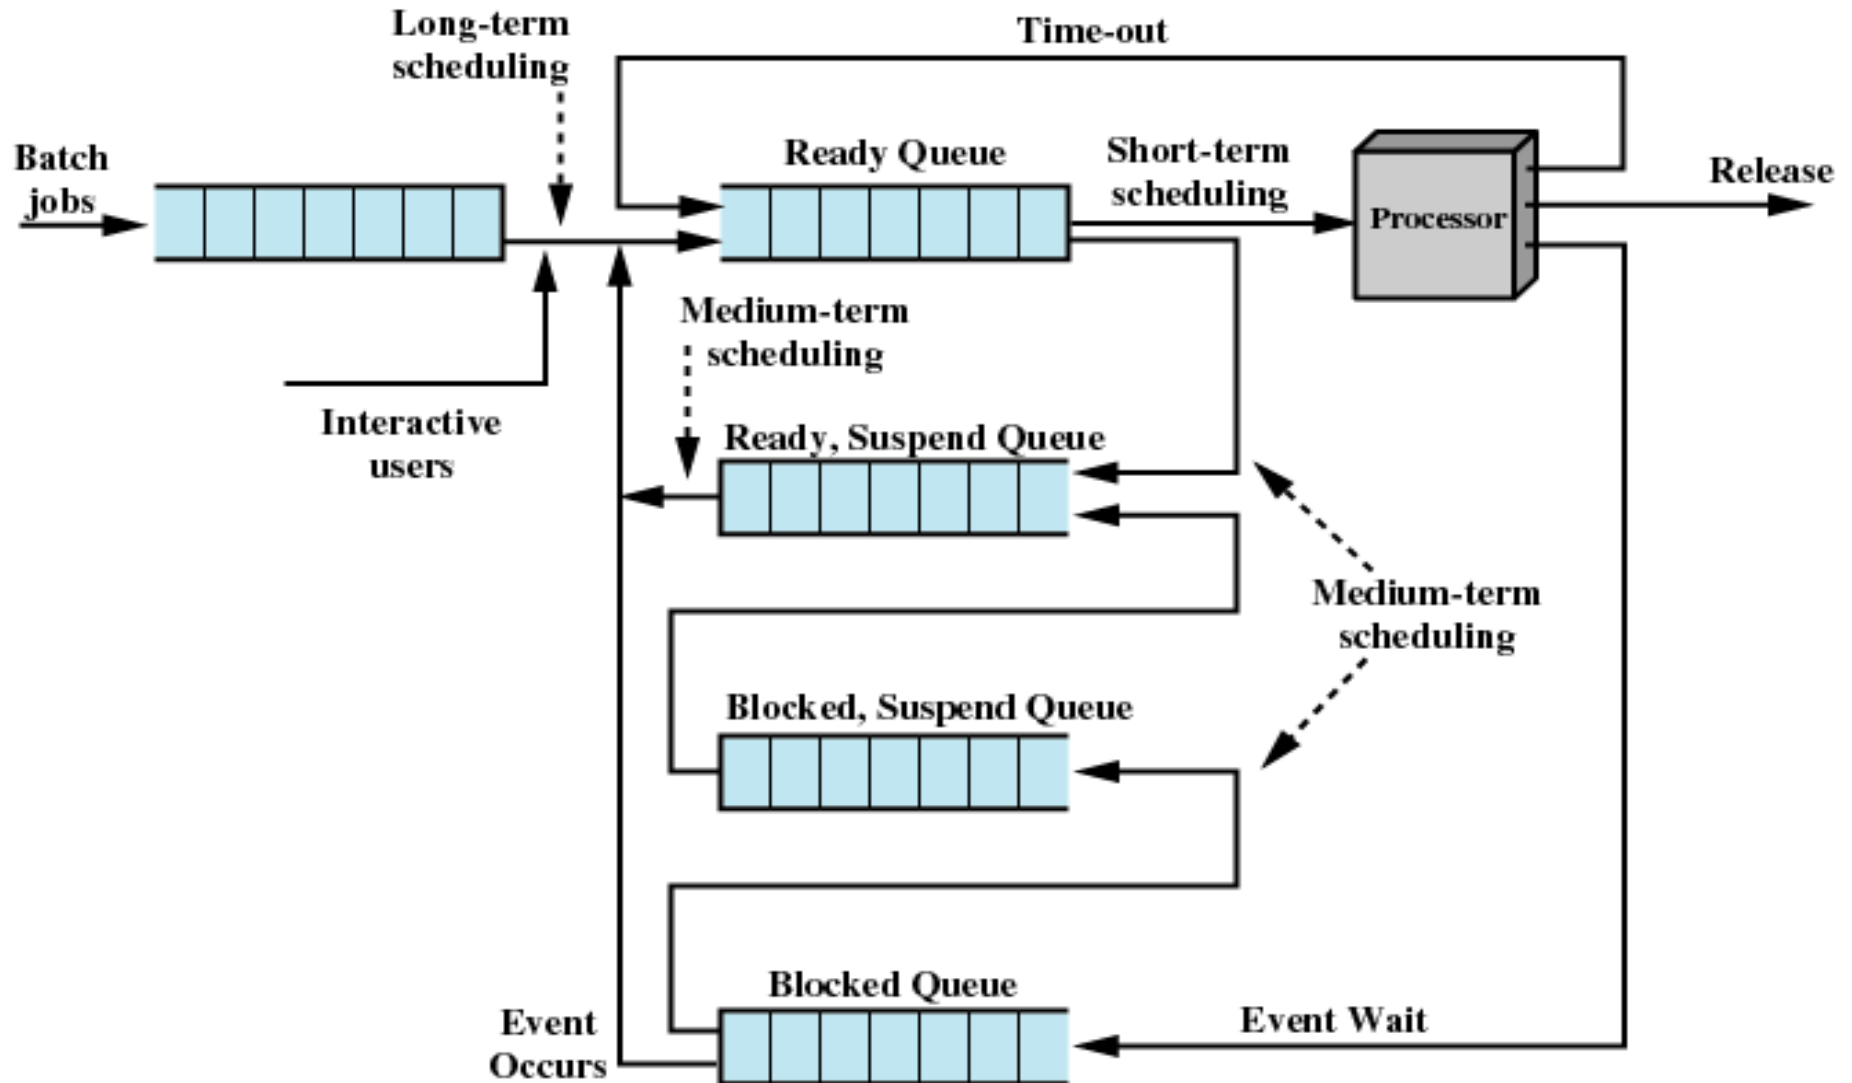
- Which I/O request (of which process) to dispatch to I/O device for handling
- Consider state of external device



- User-oriented:
  - Response time: elapsed time between submission of a request until there is output
  - Interactivity: user *perceives* system as “responsive”
- System-oriented (hardware and resources):
  - Effective and efficient utilization of processor
- Performance-related:
  - Quantitative / measurable properties
  - Examples: response time, throughput
- Non-functional:
  - Algorithmic properties
  - Examples: predictability, fairness

	Performance-related	Non-functional
User-oriented	<ul style="list-style-type: none"><li>• Turnaround time</li><li>• Response time</li><li>• Deadlines</li></ul>	<ul style="list-style-type: none"><li>• Predictability</li></ul>
System-oriented	<ul style="list-style-type: none"><li>• Throughput</li><li>• Processor utilization</li></ul>	<ul style="list-style-type: none"><li>• Fairness</li><li>• Enforcing priorities</li><li>• Balancing resources</li></ul>

# Scheduler Implementation: Queuing

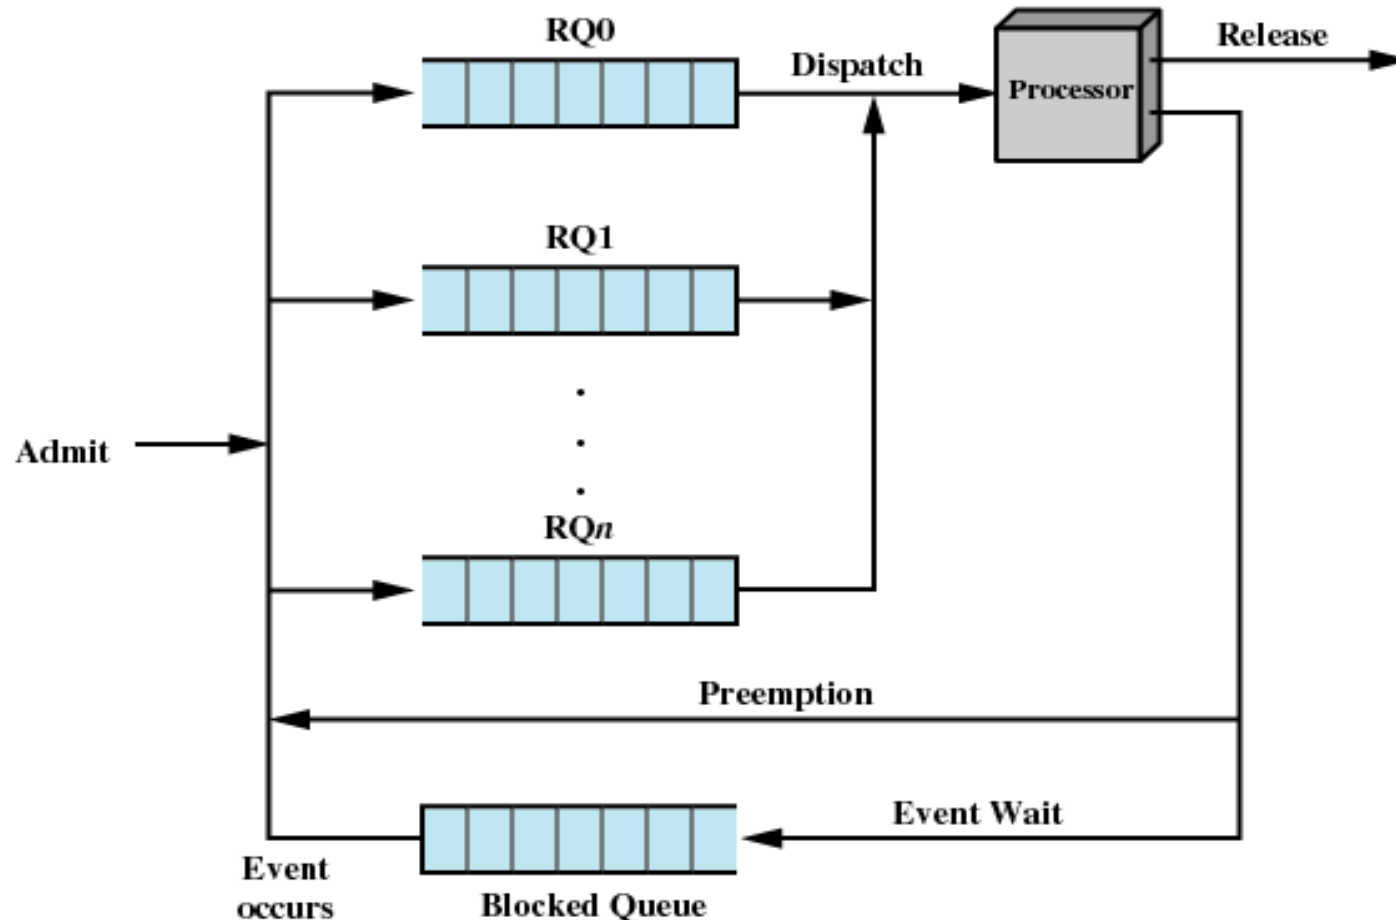


- Non-preemptive
  - Current process explicitly yields CPU
    - Cooperative multitasking, e.g., Windows (<95), Mac OS (<X)
  - Once a process is in running state, it will continue until it terminates or blocks itself for I/O
- Preemptive
  - OS may interrupt current process
    - Transparent to process
      - Preemptive multitasking, e.g., Windows ( $\geq 95$ ), Mac OS X, Unix
  - Allows for better scheduling since no process can monopolize CPU

- Some processes are more *important* than other processes, i.e., should get more CPU cycles than others
  - Similar for other resources
- Scheduling is controlled by per-process priorities
  - OS internal vs. user-visible priorities
- Scheduler will always choose a process of higher priority over one of lower priority
- Lower-priority processes may suffer starvation, i.e. are never scheduled and do not make *any* progress

# Priority Implementation: Queuing

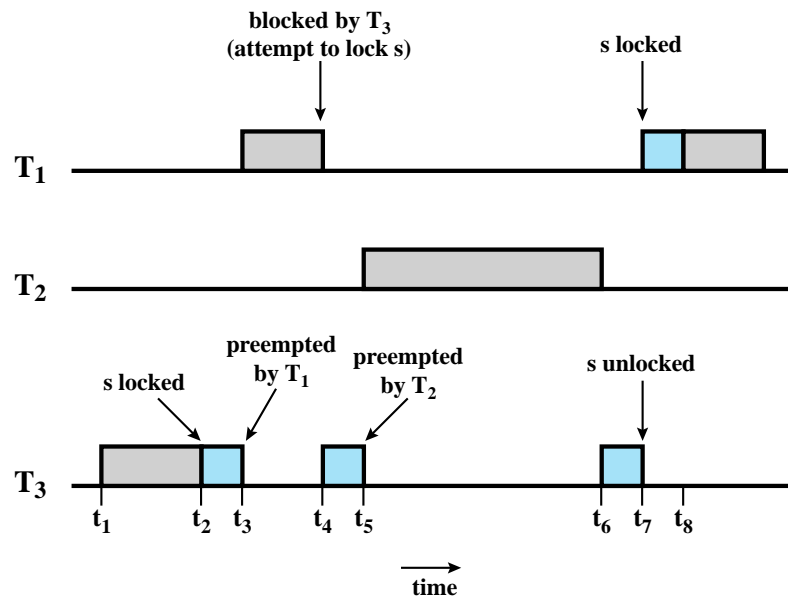
- Have multiple ready queues to represent each level of priority
- Move process data between queues according to scheduling algorithm



# Priority Inversion and Inheritance

## Problem: Priority Inversion

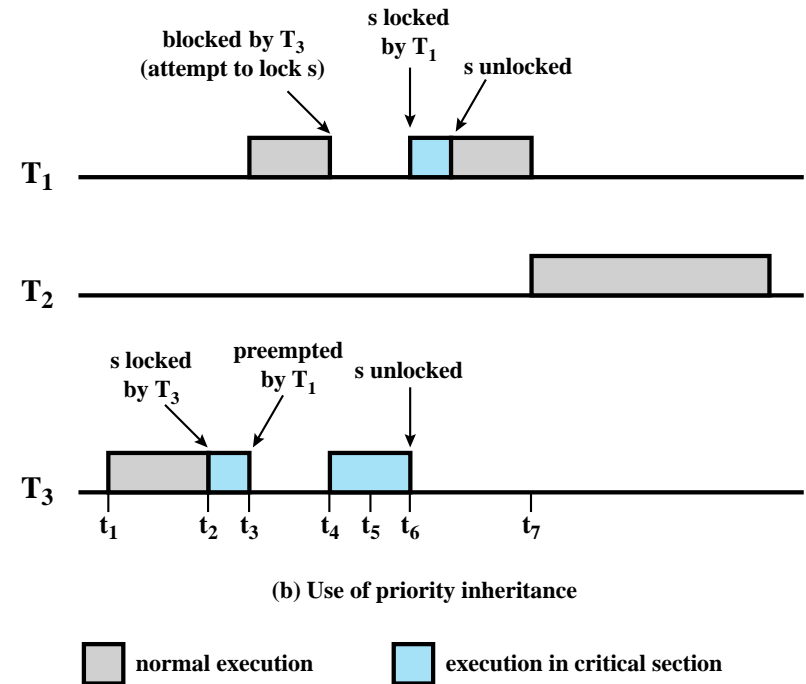
- Occurs when circumstances within the system force a higher priority task to wait for a lower priority task



(a) Unbounded priority inversion

## Solution: Priority Inheritance

- Lower-priority task inherits priority of any higher priority task pending on a resource they share



- Non-preemptive
  - First-Come-First-Served (FCFS)
  - Shortest Process Next (SPN)
  - Highest Response Ratio Next (HRRN)
- Preemptive
  - Shortest Remaining Time (SRT)
  - Round-Robin
  - Feedback

Example workload

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2




# First-Come-First-Served (FCFS)

- New process placed at end of Ready queue
- When current process ceases to execute, oldest process in the Ready queue is selected

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

## First-Come-First Served (FCFS)




- Short process may have to wait a very long time before it can execute
  - Poor response time / interactivity
- Favors CPU-bound processes
  - I/O processes have to wait until CPU-bound process completes, since I/O processes frequently call into OS

# Shortest Process Next (SPN)

- Process with shortest expected processing time is selected
  - OS may abort processes with incorrect time estimates
- Short processes jump ahead of longer processes

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

**Shortest Process  
Next (SPN)**



- Improves interactivity (based on assumption that short processes are due to user interaction)
- Predictability of longer processes is reduced
- Possibility of starvation for longer processes


# Highest Response Ratio Next (HRRN)

- Choose next process with the highest ratio

$$\frac{\text{time spent waiting} + \text{expected service time}}{\text{expected service time}}$$

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

**Highest Response Ratio Next (HRRN)**




- Even long process will run eventually
- Generally, predictable response times not feasible without preemption

# Shortest Remaining Time (SRT)

- Ready queue is sorted by remaining processing time
  - Requires estimate of remaining processing time
- New processes may preempt current process upon arrival
  - Preemptive version of shortest process next policy

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

**Shortest Remaining Time (SRT)**




- Improved response time of short processes by using preemption
  - Limited additional overhead due to process switches upon process creation
- But happens to interactive requests that don't spawn a new process?

# Round-Robin

- Each process may use CPU for given amount of time
  - Process preemption based on clock interrupt generated at intervals, i.e., time slicing
  - Time quantum  $q$  as tunable parameter
- When interrupt occurs, currently running process is placed in Ready queue, next ready job is selected

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2




- Initial support for interactivity
- Scheduling overhead (scheduling decision, process switch)
  - Tradeoff between interactivity and efficiency, directly tunable by  $q$
- Problematic for I/O processes that hardly ever use full quantum

- Processes start in the queue with highest priority RQ0 and move to queues with lower priority after each time slice
  - Multiple queues with different priorities
- For fairness, allow longer time slices  $q$  for queues RQ $i$

**Feedback**  
 $q = 1$

**Feedback**  
 $q = 2^i$

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2



- Penalize long running processes
- No need to know remaining execution time of process

# Qualitative Comparison of Policies

	Selection Function	Decision Mode	Throughput	Response Time	Overhead	Effect on Processes	Starvation
<b>FCFS</b>	$\max[w]$	Nonpreemptive	Not emphasized	May be high, especially if there is a large variance in process execution times	Minimum	Penalizes short processes; penalizes I/O bound processes	No
<b>Round Robin</b>	constant	Preemptive (at time quantum)	May be low if quantum is too small	Provides good response time for short processes	Minimum	Fair treatment	No
<b>SPN</b>	$\min[s]$	Nonpreemptive	High	Provides good response time for short processes	Can be high	Penalizes long processes	Possible
<b>SRT</b>	$\min[s - e]$	Preemptive (at arrival)	High	Provides good response time	Can be high	Penalizes long processes	Possible
<b>HRRN</b>	$\max\left(\frac{w + s}{s}\right)$	Nonpreemptive	High	Provides good response time	Can be high	Good balance	No
<b>Feedback</b>	(see text)	Preemptive (at time quantum)	Not emphasized	Not emphasized	Can be high	May favor I/O bound processes	Possible

$w$  = time spent waiting,  $e$  = time spent in execution so far,  $s$  = total service time required by process, including  $e$

# Quantitative Comparison of Policies

	Process	A	B	C	D	E	
	Arrival Time	0	2	4	6	8	
	Service Time ( $T_S$ )	3	6	4	5	2	Mean
FCFS	Finish Time	3	9	13	18	20	
	Turnaround Time ( $T_T$ )	3	7	9	12	12	8.60
	$T_T/T_S$	1.00	1.17	2.25	2.40	6.00	2.56
RR $q = 1$	Finish Time	4	18	17	20	15	
	Turnaround Time ( $T_T$ )	4	16	13	14	7	10.80
	$T_T/T_S$	1.33	2.67	3.25	2.80	3.50	2.71
RR $q = 4$	Finish Time	3	17	11	20	19	
	Turnaround Time ( $T_T$ )	3	15	7	14	11	10.00
	$T_T/T_S$	1.00	2.5	1.75	2.80	5.50	2.71
SPN	Finish Time	3	9	15	20	11	
	Turnaround Time ( $T_T$ )	3	7	11	14	3	7.60
	$T_T/T_S$	1.00	1.17	2.75	2.80	1.50	1.84
SRT	Finish Time	3	15	8	20	10	
	Turnaround Time ( $T_T$ )	3	13	4	14	2	7.20
	$T_T/T_S$	1.00	2.17	1.00	2.80	1.00	1.59
HRRN	Finish Time	3	9	13	20	15	
	Turnaround Time ( $T_T$ )	3	7	9	14	7	8.00
	$T_T/T_S$	1.00	1.17	2.25	2.80	3.5	2.14
FB $q = 1$	Finish Time	4	20	16	19	11	
	Turnaround Time ( $T_T$ )	4	18	12	13	3	10.00
	$T_T/T_S$	1.33	3.00	3.00	2.60	1.5	2.29
FB $q = 2^i$	Finish Time	4	17	18	20	14	
	Turnaround Time ( $T_T$ )	4	15	14	14	6	10.60
	$T_T/T_S$	1.33	2.50	3.50	2.80	3.00	2.63





# Multiprocessor and Real-Time Scheduling

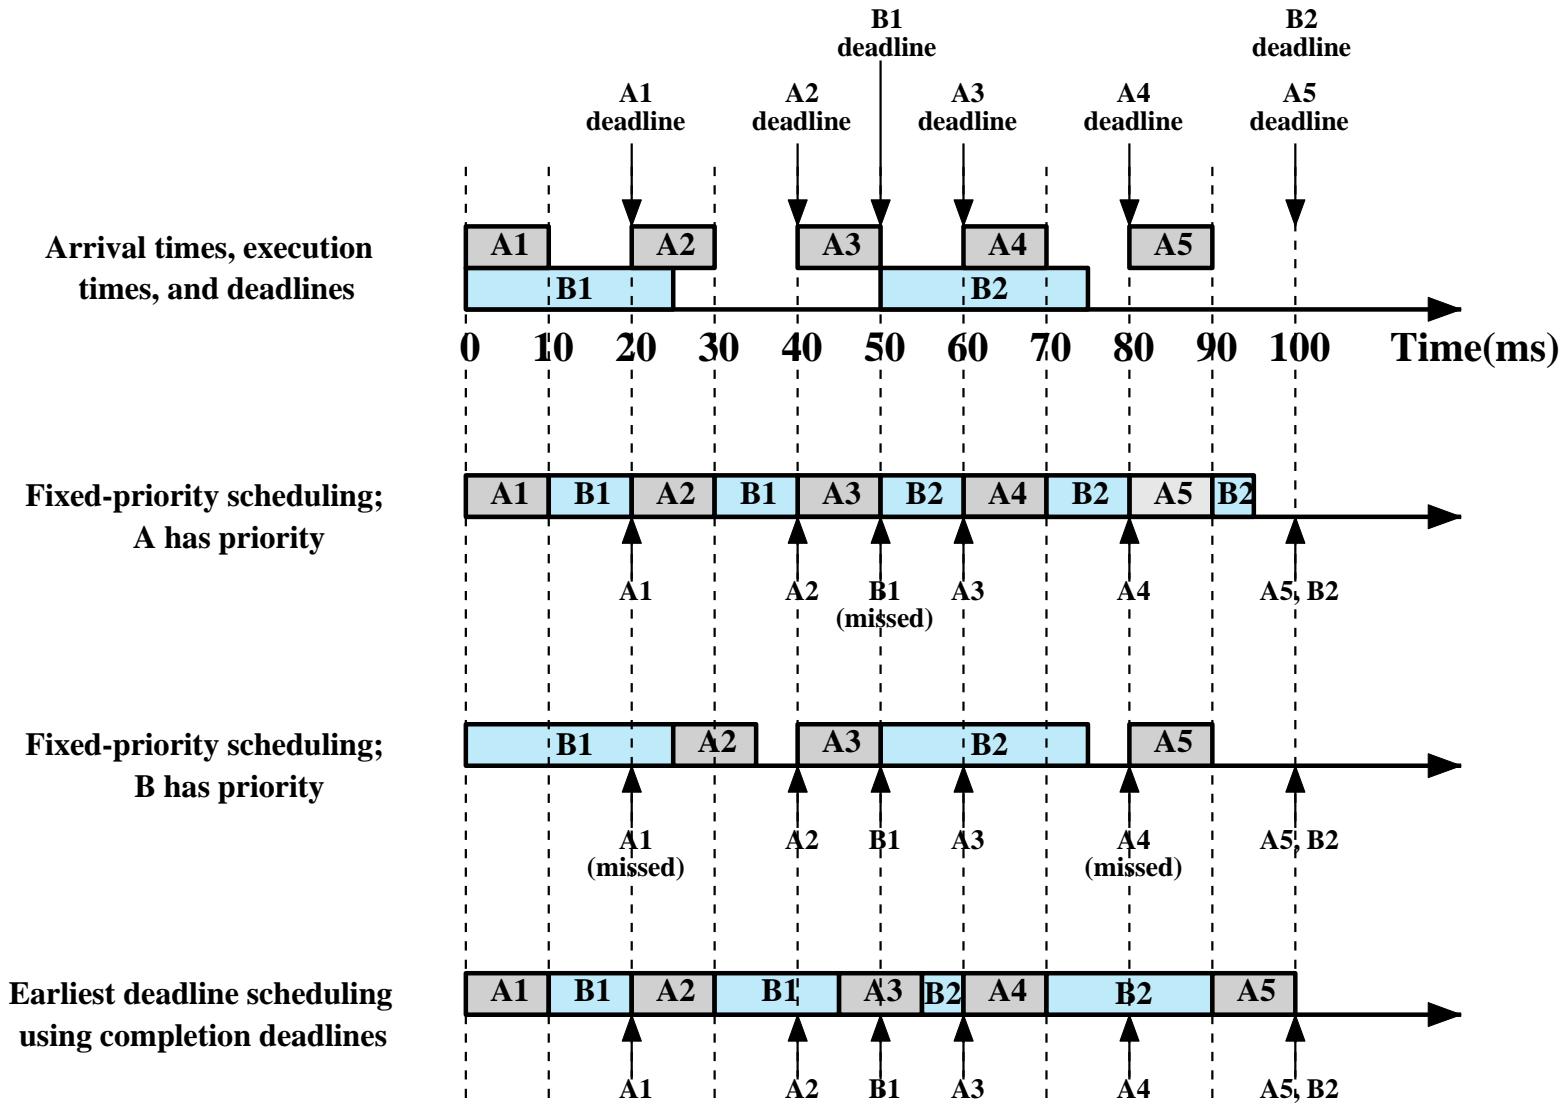
# Multiprocessor Scheduling

- Assignment of processes to processors
  - Permanently assign process to a processor
  - Treat processors as a pooled resource and assign process to processors on demand
    - Possibly move running process between processors (expensive!)
- Architectures
  - Global queue: schedule to any available processor
  - Master/slave: Key kernel functions always run on particular processor, master is responsible for scheduling
  - Peer: Operating system can execute on any processor, each processor does self-scheduling
- Use of multiprogramming on individual processors
- Actual dispatching of processes

# Real-Time Scheduling

- Correctness of system depends
  - on logical result of the computation
  - **AND** on time at which the results are produced
- Tasks or processes attempt to control or react to events that take place in outside world
- Examples:
  - Control of laboratory experiments
  - Process control in industrial plants
  - Robotics
  - Air traffic control
  - Telecommunications
  - Military command and control systems
- Real-time applications are *not concerned with speed* but with *completing tasks*

# Real-Time Scheduling: Examples





# Examples

- Multilevel feedback using round robin within each priority queue
- If running process does not block or complete within one second, it is preempted
- Priorities are recomputed once per second
- Base priority (set upon process creation) divides all processes into fixed bands of priority levels

# Examples: UNIX SVR4 Scheduling

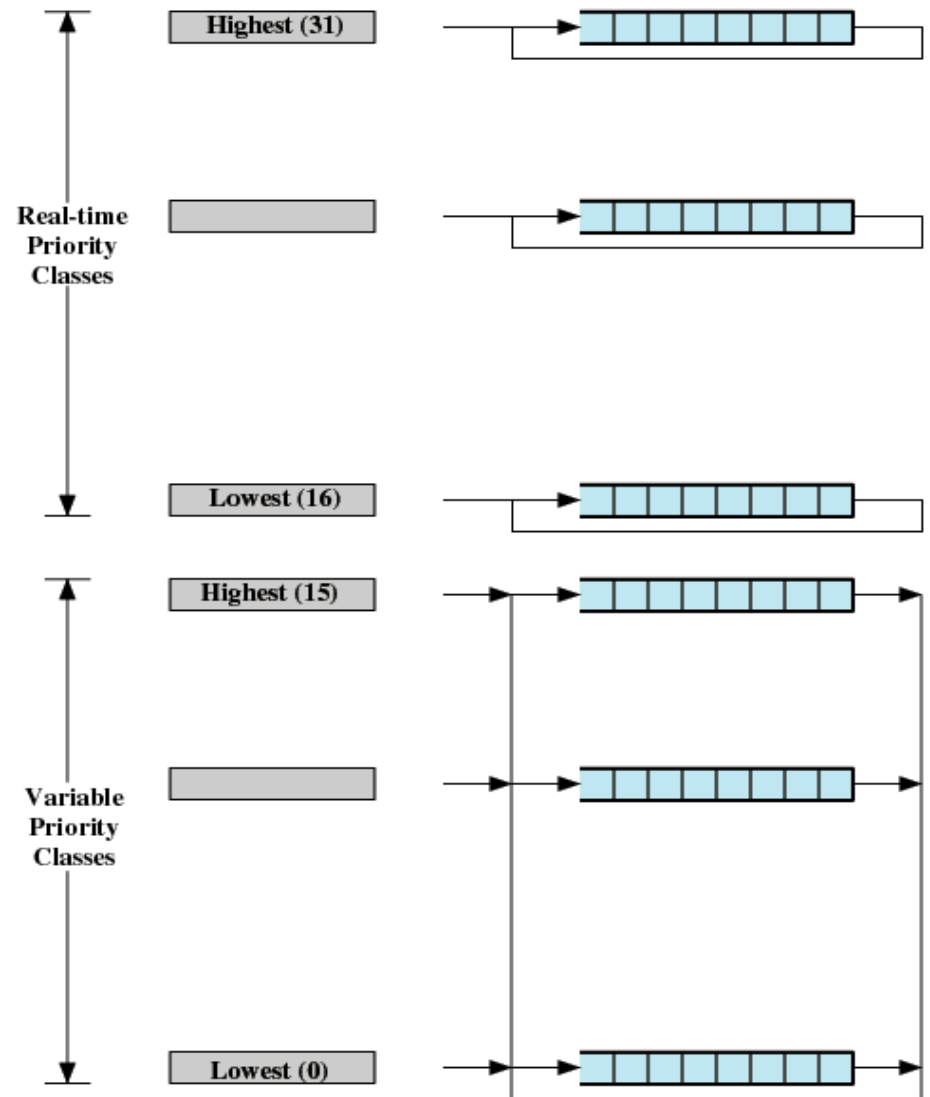

- Preemptable static priority scheduler
- Introduces set of 160 priority levels divided into three priority classes
  - Highest preference to real-time processes
  - Next-highest to kernel-mode processes
  - Lowest preference to other user-mode processes
- In-kernel preemption points, i.e. long running kernel operations may be preempted

- SVR4 Priority Classes:
  - Real time (159 – 100)
  - Kernel (99 – 60)
  - Time-shared (59-0)

Priority Class	Global Value	Scheduling Sequence
Real-time	159	first ↓
	.	
	.	
	.	
	100	
Kernel	99	
	.	
	60	
Time-shared	59	↓ last
	.	
	.	
	.	
	0	

# Examples: Windows Scheduling

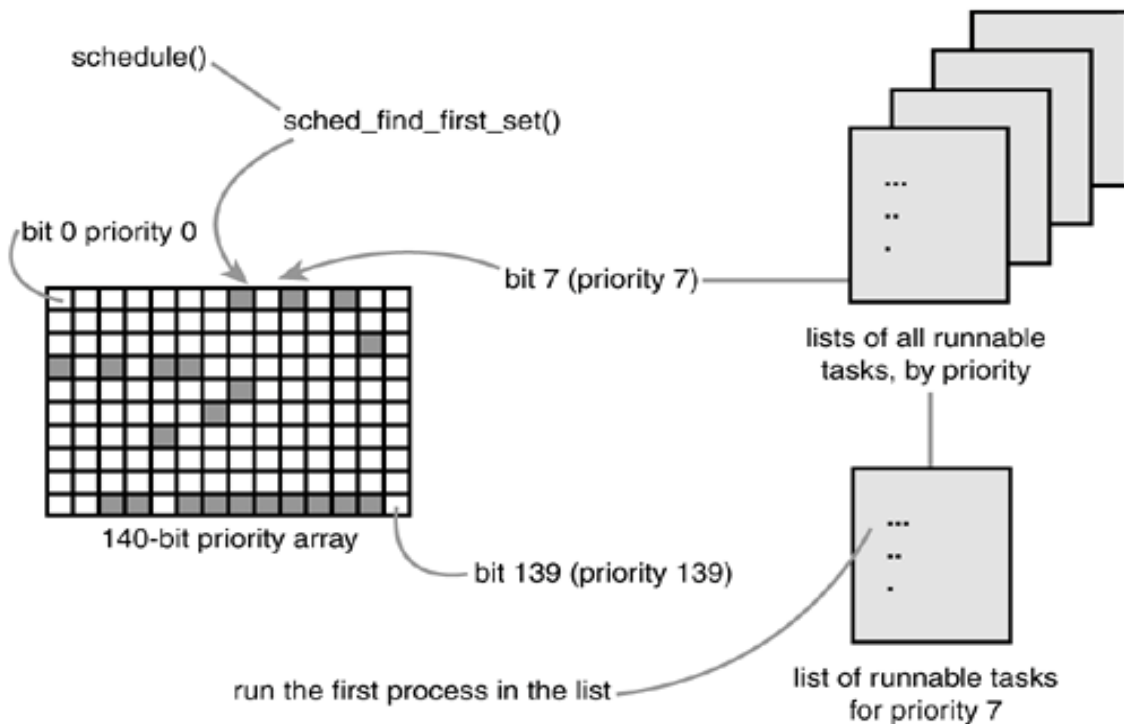
- Priorities organized into two bands or classes
  - Real time
  - Variable
- Priority-driven preemptive scheduler within each class





# Example: Linux O(1) Scheduling

- Scheduling algorithm needs to scale with number of processes
  - Variable overhead unacceptable for real-time systems



## ➤ Linux O(1) scheduler

- Active/expired bit arrays for priorities; one list per priority
  - Priority assigned based on
    - Static (process) priority
    - Heuristics to determine interactivity requirements, e.g. CPU- vs. I/O-bound
  - Process timeslice (i.e. runtime in relation to other processes) calculated when process moves from active to expired state
  - Switch from active to expired bit array when all processes have used their timeslice
- Scheduling decision in constant time

# Related System Calls

- `int sched_yield(void)`
  - Voluntarily yield processor, e.g. when waiting for input
- `int getpriority(int which, int who)`
- `int setpriority(int which, int who, int prio)`
  - Get/set priority of user, group or process (**which**) with ID **who**
  - Library interface: `int nice(int inc)`
    - Increment how nice you are; only root is allow not to be nice
- `int sched_get_priority_max(int policy)`
- `int sched_get_priority_min(int policy)`
  - Returns max/min priority values for given scheduling **policy**

## Related System Calls (Linux, cont.)

- `int sched_setscheduler(pid_t pid, int policy, const struct sched_param *param)`
- `int sched_getscheduler(pid_t pid)`
  - Controls which scheduling policy to use for a process
  - Policies are `SCHED_BATCH`, `SCHED_FIFO`, `SCHED_RR` and `SCHED_OTHER`
- `int sched_setparam(pid_t pid, const struct sched_param *param)`
- `int sched_getparam(pid_t pid, struct sched_param *param)`
  - Get/set policy specific scheduling parameters
- `int sched_setaffinity(pid_t pid, unsigned int cpusetsize, cpu_set_t *mask)`
- `int sched_getaffinity(pid_t pid, unsigned int cpusetsize, cpu_set_t *mask)`
  - Controls on which CPU in multi-processor system a process can/should run

1. Introduction and Motivation
2. Subsystems, Interrupts and System Calls
3. Processes
4. Memory
- 5. Scheduling**
6. I/O and File System
7. Booting, Services, and Security