


# Operating Systems & Computer Networks

## Host-to-Network II

- Data Link Layer
- Framing, Flow Control
- Error Detection / Correction
- Point-to-Point Protocol



8. Networked Computer & Internet

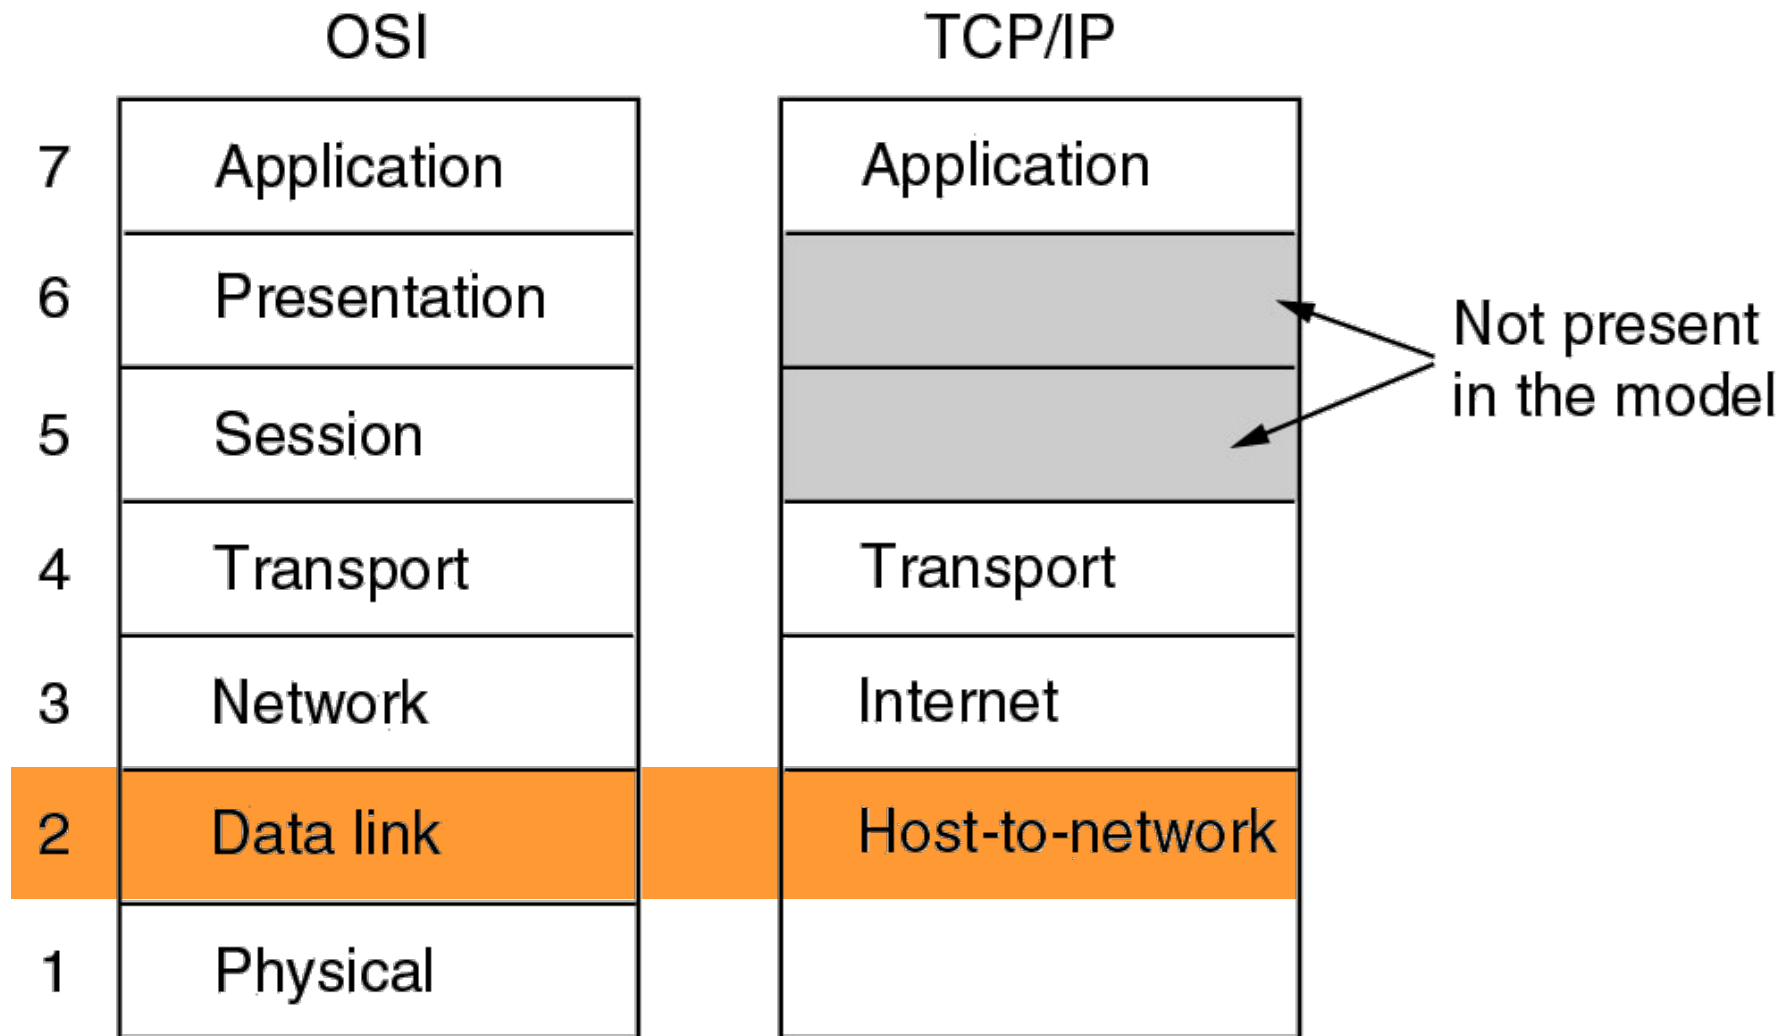
9. Host-to-Network I

**10. Host-to-Network II**

11. Host-to-Network III

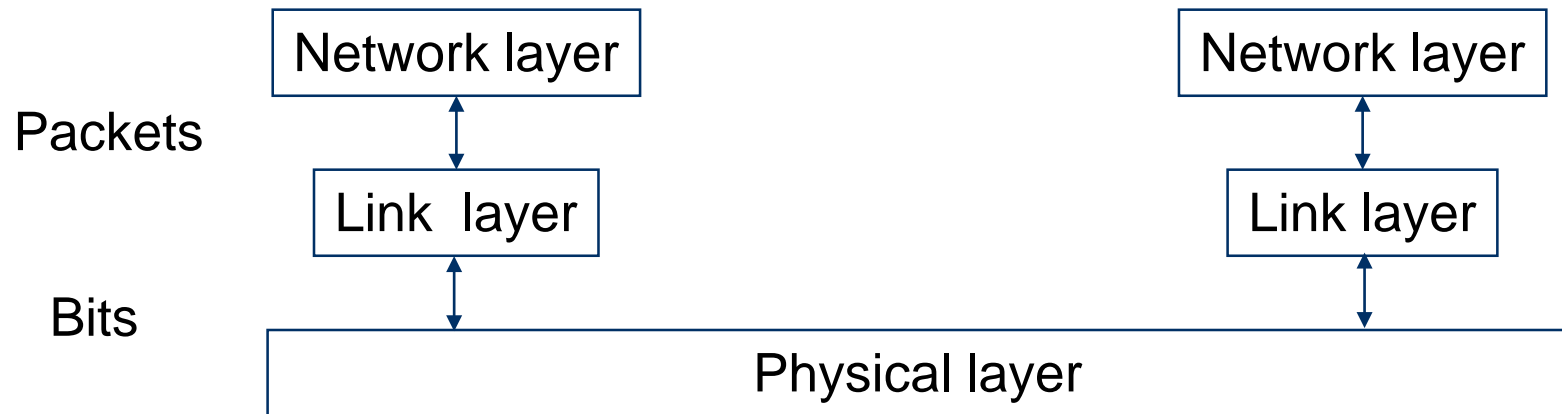
12. Internetworking

13. Transport Layer



# Setting for Data Link Layer

- Link layer sits on top of physical layer
  - Can thus use a bit stream transmission service
  - But: Service might have incorrect bits
- Expectations of the higher layer (networking layer)
  - Wants to use either a packet service
    - Rarely, a bit stream service
  - Does not really want to be bothered by errors
  - Does not really want to care about issues at the other end



# Services of Data Link Layer

Given setting and goals, the following services are required:

- Transparent communication between two directly connected nodes
- Framing of a physical bit stream into a structure of frames/packets
  - Frames can be retransmitted, scheduled, ordered, ...
- Error control
  - Detection and correction
- Connection setup and release
  - Signaling and resource management on hosts
- Acknowledgement-based protocols
  - Make sure that a frame has been transmitted
- Flow control
  - Arrange for appropriate transmission speed between hosts

# Framing

# Link Layer Functions – Framing

- How to turn bit stream abstraction (as provided by physical layer) into individual, well demarcated *frames*?
  - Usually necessary to provide error control
    - Not obvious how to do that over a bit stream abstraction
  - Frames and packets are really the same thing
    - Term “frames” used in link layer context by convention
- Additionally: Fragmentation and reassembly
  - If network layer packets are longer than link layer packets
  - Commonly network layer set Maximum Transmission Unit (MTU) to a value that avoids link layer fragmentation
    - 1492 bytes on 802.3 (Ethernet)
    - 2272 bytes on 802.11 (WLAN)

- How to turn a bit stream into a sequence of frames?
  - More precisely: how does a receiver know when a frame starts and when it ends?

Delivered by physical layer:

0110010101110101110010100010101010101010101100010



Start of frame (?)



End of frame (?)


- Note: Physical layer might try to detect and deliver bits when the sender is not actually transmitting anything
  - Receiver tries to get any information from the physical medium



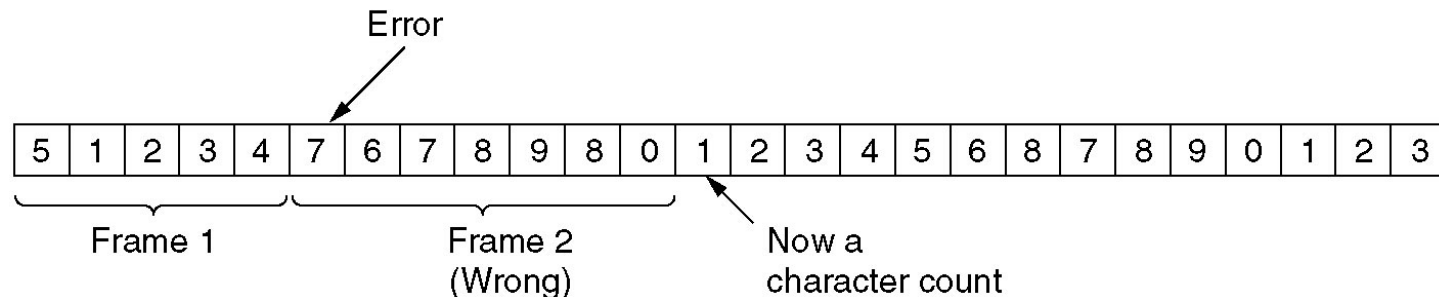
# Framing by Character / Byte Count

- Idea: Announce number of bits (bytes, characters) in a frame to the receiver

➤ Put this information at beginning of a frame – *frame header*



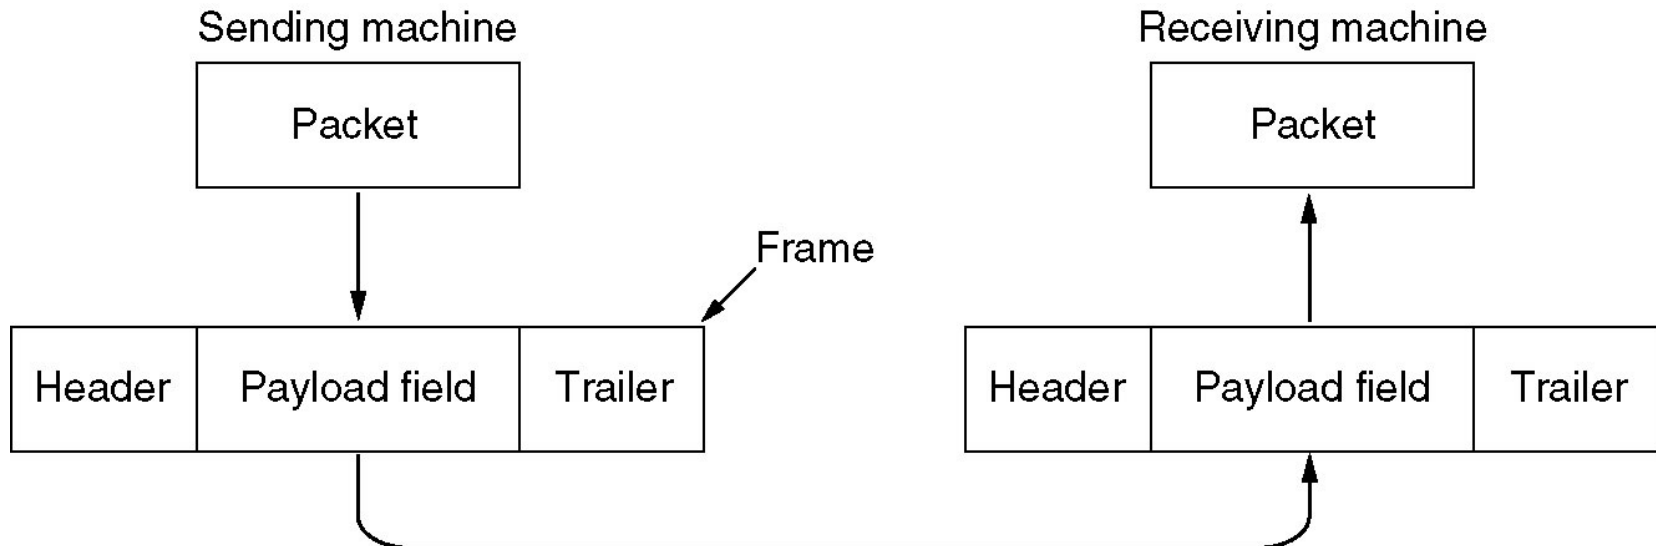
- Problem: What happens if *count* information itself is damaged during transmission?
  - Receiver will lose frame synchronization and produce different sequence of frames than original one





# Basic Technique: Control Header

- Albeit “character count” is not a good framing technique, it illustrates an important technique: *headers*
  - If sender has to communicate administrative or control data to receiver, it can be added to actual packet content (“payload”)
  - Usually at the start of the packet; sometimes at the end (“trailer”)
  - Receiver uses headers to learn about sender’s intention
- Same principle applicable to all packet-switched communication



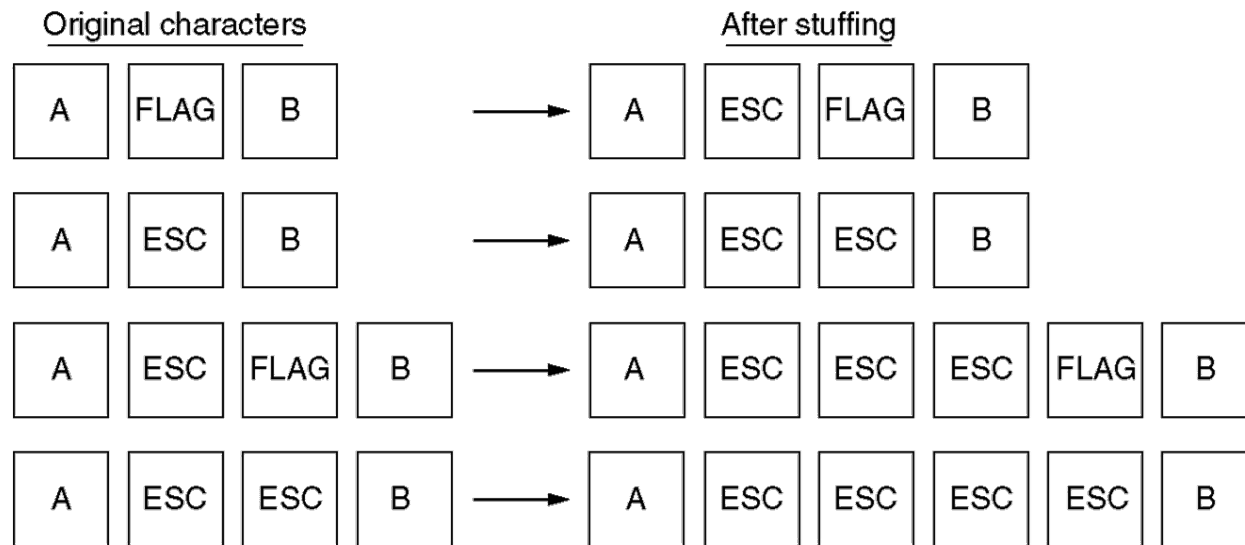


# Framing by Flag Bytes / Byte Stuffing

- Use dedicated flag bytes to demarcate start/stop of frame



- What happens if the flag byte appears in the payload?
- Escape it with a special control character – *byte stuffing*
  - If that appears, escape it as well



# Framing by Flag Bit Patterns / Bit Stuffing

- Byte stuffing is closely tied to characters/bytes as fundamental unit – often not appropriate
- Use same idea, but stick with the bit stream abstraction of the physical layer
  - Use bit pattern instead of flag byte – often, 01111110
    - Actually, it IS a flag byte
  - Bit stuffing process:
    - Whenever sender sends five 1s in a row, it automatically adds a 0 into bit stream – except in flag pattern
    - Receiver throws away (“destuffs”) any 0 after five 1s

Original payload      (a) 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0

After bit stuffing      (b) 0 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 1 0



Stuffed bits

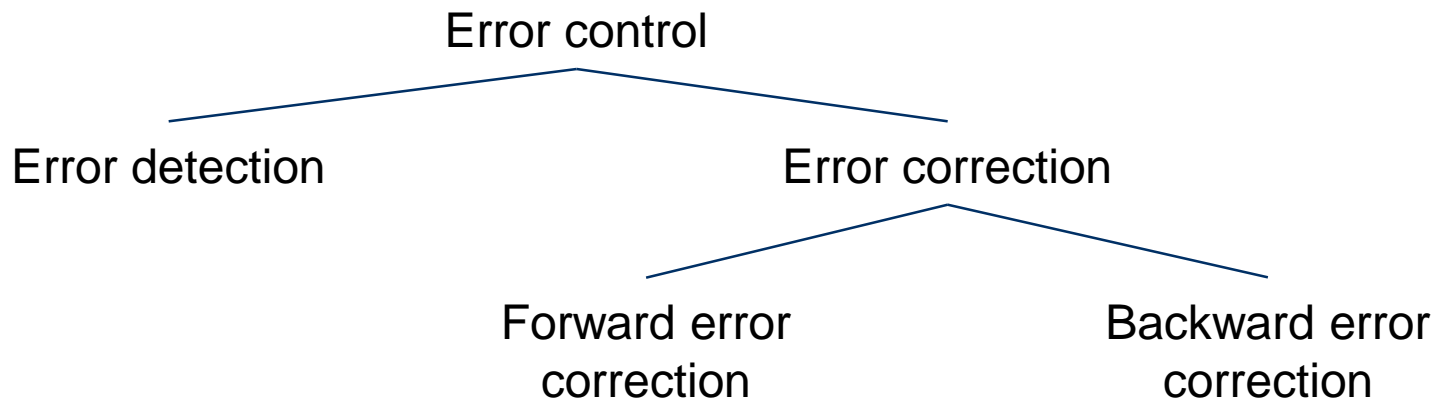
After destuffing      (c) 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0

# Framing by Coding Violations

- Suppose the physical layer's encoding rules "bits ! signals" still provide some options to play with
  - Not all possible combinations that physical layer can represent are used to represent bit patterns
    - Example: Manchester encoding: only low/high and high/low is used
- When "violating" these encoding rules, data can be transmitted, e.g., start and end of frame
  - Example: Manchester – use high/high or low/low
    - This drops self-clocking feature of Manchester, but clock synchronization is sufficiently good to hold for a short while
- Powerful and simple scheme, e.g. used by Ethernet networks
  - But raises questions regarding bandwidth efficiency as coding is *obviously* not optimal

# Error Control

- Error detection – Check for incorrect bits
- Error correction – Correct erroneous bits
  - Forward error correction (FEC) – invest effort *before* error happened; avoid delays in dealing with it
    - Redundancy / overhead
  - Backward error correction – invest effort *after* error happened; try to repair it ➔ ARQ (Automatic Repeat reQuest)
    - Delays



➤ Usually build on top of framing



# Error Detection: Cyclic Redundancy Check (CRC)

- CRC can check arbitrary, unstructured sequence of bits
- A check sequence (CRC code) is appended to checked data
  - Typically calculated by a feedback shift register in hardware

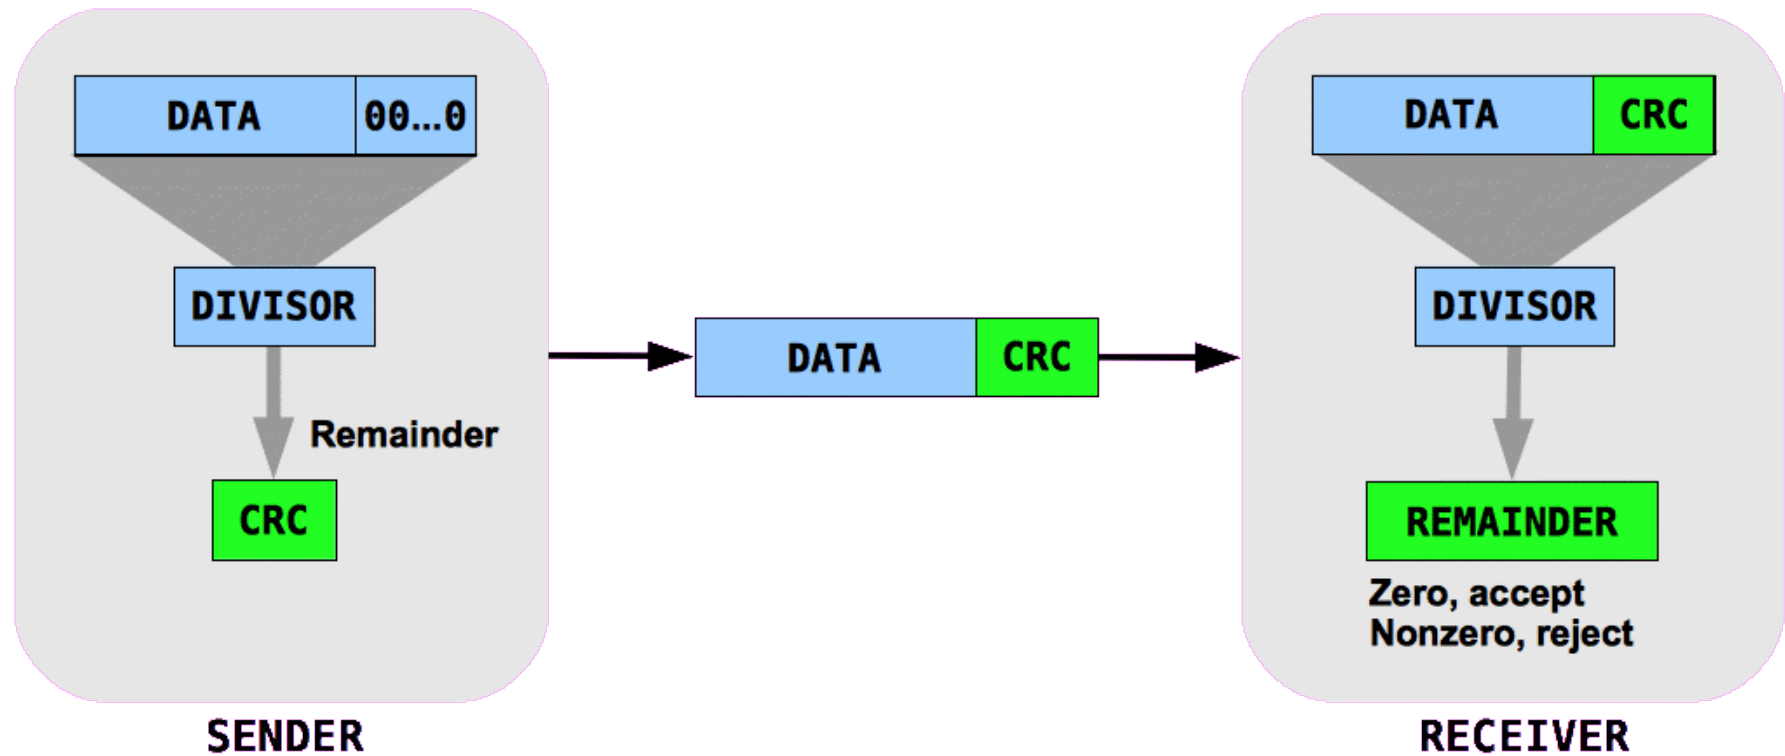


- When calculating the CRC code a generator polynomial is used which is known to both sender and receiver
- Calculation of CRC code
  1. View bit sequence as polynomial with binary coefficients:
    - **100010** is viewed as  $1 \cdot x^5 + 0 \cdot x^4 + 0 \cdot x^3 + 0 \cdot x^2 + 1 \cdot x^1 + 0 \cdot x^0$
  2. Expand polynomial with n 0s, n is the degree of the generator polynomial
  3. Divide expanded bit sequence (i.e. polynomial) by generator polynomial
    - CRC code is the remainder of the division, result is discarded
  4. Receiver again divides received bit sequence (including the CRC code) by generator polynomial
    - If no error occurred the remainder is 0





# Illustration of CRC





## CRC Example (Sender)

- Transmitted payload: 110011
- Generator polynomial:  $x^4 + x^3 + 1$ 
  - Translates into sequence of coefficients: 11001
    - Addition or subtraction equal simple bitwise XOR
      - Special arithmetic for polynomials modulo 2
- Length of CRC = degree of generator polynomial = 4
- Calculation of CRC:

$$110011\textcolor{red}{0000} \div 11001 = 100001 \pmod{2}$$

$$\begin{array}{r} 11001 \\ \hline \end{array}$$

$$0000010000$$

$$\begin{array}{r} 11001 \\ \hline \end{array}$$

$$\textcolor{blue}{1001} = \text{remainder}$$

- Transmitted bit sequence: 110011 $\textcolor{blue}{1001}$



# CRC Example (Receiver)

- Reception of a correct bit sequence:

$$1100111001 \div 11001 = 100001 \pmod{2}$$

$$\begin{array}{r} 11001 \\ 0000011001 \\ \underline{11001} \\ 00000 = \text{remainder} \end{array}$$

- No remainder, thus the received bits *should* be error free

- Reception of a erroneous bit sequence:

$$1111111000 \div 11001 = 101001 \pmod{2}$$

$$\begin{array}{r} 11001 \\ 0011011 \\ \underline{11001} \\ 00010000 \\ \underline{11001} \\ 01001 = \text{remainder} \neq 0 \end{array}$$

- There is a remainder unequal 0, thus there was *definitely* a transmission error



# Properties of CRC

## CRC can detect the following errors:

- All single bit errors
- All double bit errors (if  $(x^k + 1)$  is not divisible by generator polynomial for  $k \leq \text{frame length}$ )
- All errors affecting an odd number of bits (if  $(x+1)$  is a factor of the generator polynomial)
- All error bursts of length  $\leq$  degree of generator polynomial

## Internationally standardized generator polynomials:

- CRC-12  $= x^{12} + x^{11} + x^3 + x^2 + x + 1$
- CRC-16  $= x^{16} + x^{15} + x^2 + 1$
- CRC-CCITT  $= x^{16} + x^{12} + x^5 + 1$

## CRC-16 and CRC-CCITT detect

- All single and double errors
- All errors affecting an odd number of bits
- All error bursts of length  $\leq 16$
- 99,997% of all error bursts of length 17
- 99,998% of all error bursts of length  $\geq 18$



# FEC Example (Sender)

- CRC uses redundancy to *detect* errors
- FEC uses redundancy to *correct* errors
- Simple FEC scheme: Repeat data several times, then use majority decision
  - Problematic with regard to overhead
- More advanced: XOR, Reed-Solomon, ...
- Example (XOR):
  - To send the following three packets:  
0101 - P1  
1111 - P2  
0000 - P3
  - 1. Calculate a fourth packet using XOR:  
1010 - P4
  - 2. Send all four packets to the receiver
    - Overhead:  $(4 - 3) / 4 = 25\%$




# FEC Example (Receiver)

- For the receiver, it is sufficient to receive any three out of the four packets.
- Reconstruct missing packet based on received packets:
  - P1 is lost:  
1111  
0000  
1010  
0101 -> P1
  - P2 is lost:  
0101  
0000  
1010  
1111 -> P2
  - P3 is lost:  
0101  
1111  
1010  
0000 -> P3
- However, receiver still has to know which packet was lost
  - Requires packet numbering

# Flow Control

# Link Layer Functions – Flow Control

- Assumptions in an ideal world:
  - Sender/receiver are always ready to send/receive
  - Receiver can handle amount of incoming data
  - No errors occur that cannot be handled by FEC



- What happens if packets are lost?
- What happens if the sender floods the receiver?
  - Imagine a web server sending data to a mobile phone...




# Very Simple Solution: Stop-and-Wait

- Concentrate on one single packet
- Receiver acknowledges correct reception of the packet
- Sender has to wait for that acknowledgement before continuing with next packet

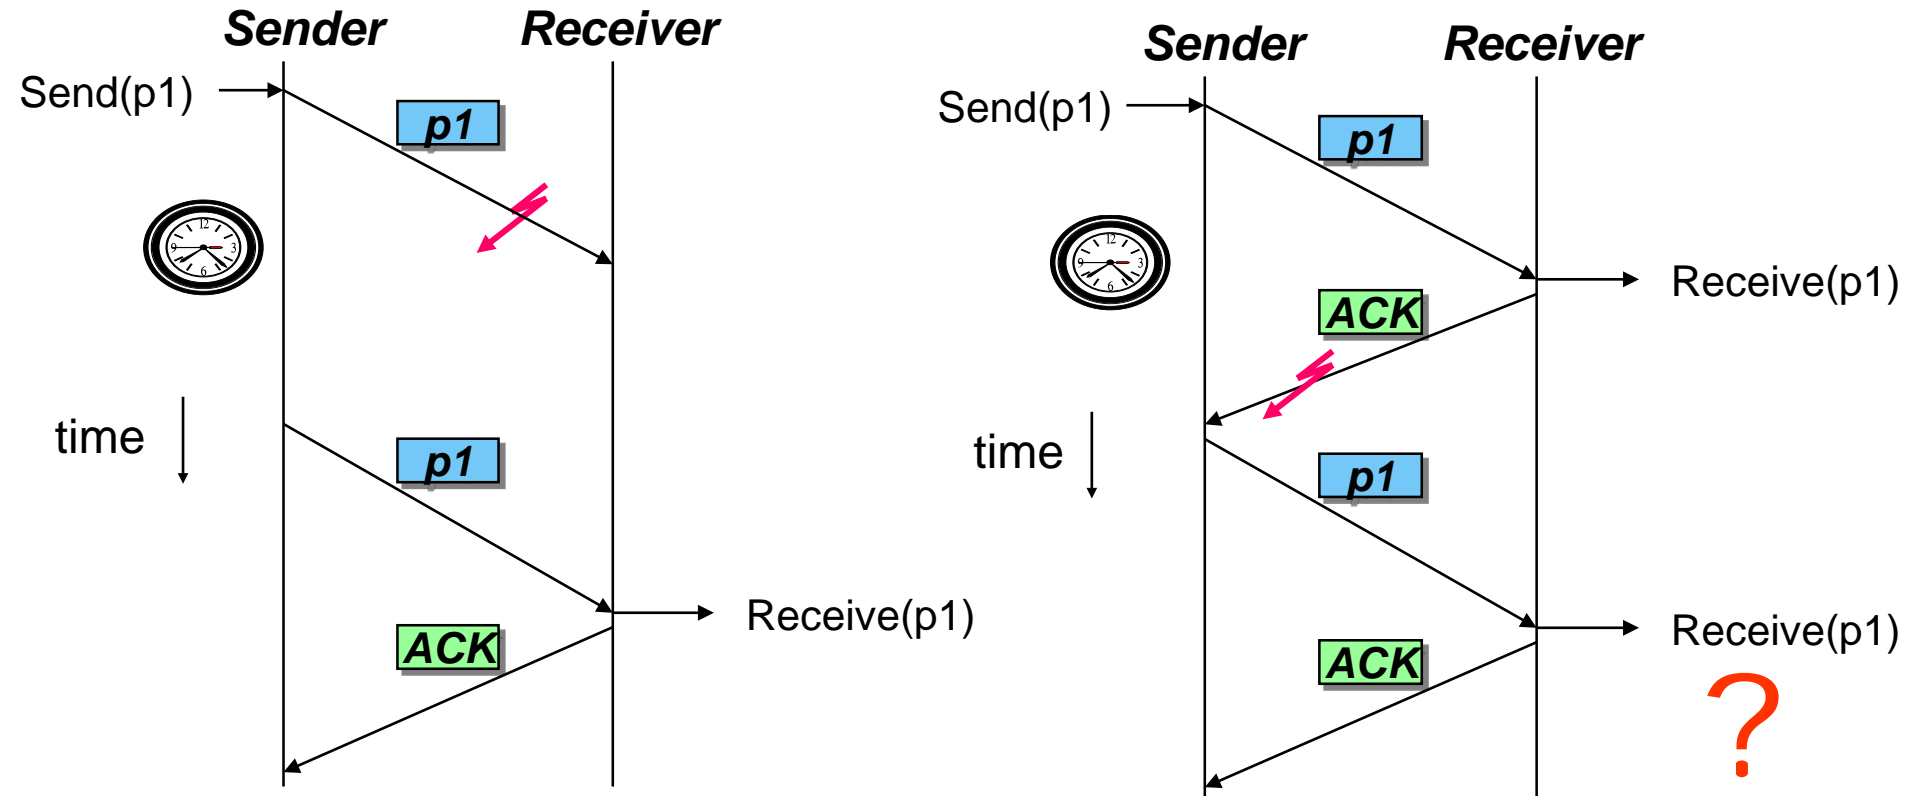
- No overloading of the receiver possible!

➤ Basic flow control



# Problems of Stop-and-Wait

- What happens if errors occur?
  - Lost packet? Lost acknowledgement? Is there a difference?
- Basic solution: ARQ (Automatic Repeat reQuest)

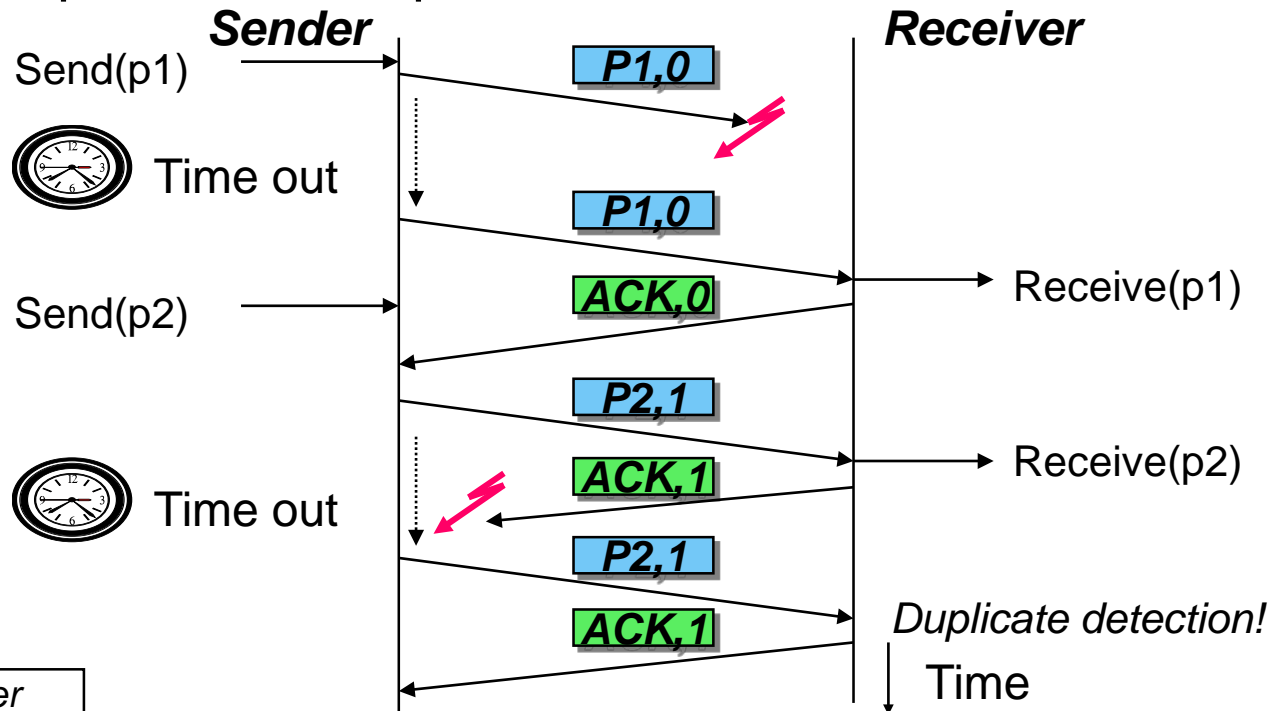


# Problem of Stop-and-Wait ARQ

- Sender cannot distinguish between lost packet and lost acknowledgement → Has to re-send the packet
  - Receiver cannot distinguish between new packet and redundant copy of old packet → Additional information is needed
- Put a *sequence number* in each packet, telling the receiver which packet it is
- Sequence numbers as *header information* in each packet
  - Simplest sequence number: 0 or 1
- Needed in packet and acknowledgement
  - One convention: In ACK, send sequence number of last correctly received packet
  - Also possible: Send sequence number of next expected packet
- Be aware: Some protocols count bytes instead of packets


# Alternating Bit Protocol

- Simple, but reliable protocol over noisy channels
- Uses 0 and 1 as sequence numbers, ARQ for retransmission
- Simple form of flow control (here combined with error control, other protocols separate these functions)



# Alternating Bit Protocol – Efficiency

- Efficiency  $\eta$ :
  - Depends on delay and bandwidth
  - Defined as ratio of time during which sender sends new information
    - assuming error-free channel in simplest case (error-considerations make efficiency discussions difficult)
  - $\eta = T_{\text{packet}} / (T_{\text{packet}} + d + T_{\text{ack}} + d)$
- Efficiency of simple alternating bit protocol is low when delay is large compared to data rate
  - Bandwidth-delay product, i.e. data in transit, not used optimally




# Improving Efficiency

## – Have More “Outstanding” Packets

- Inefficiency of alternating bit in large bandwidth-delay situations is owing to not exploiting “space” between packet and acknowledgement
- Always sending packets results in high efficiency
  - More packets are “outstanding” = sent, but not yet acknowledged
    - “Pipelining” of packets
- Not feasible with a single bit as sequence number
  - Need larger sequence number space
    - Also needs full-duplex support


Sender is  
always busy,  
efficiency  
is high



- Introduce a larger sequence number space
  - E.g.,  $n$  bits or  $2^n$  sequence numbers
- Not all of them may be allowed to be used simultaneously
  - Recall alternating bit case: 2 sequence numbers, but only 1 may be “in transit”
- Use *sliding windows* at both sender and receiver to handle these numbers
  - Sender: *sending window* – set of sequence numbers it is allowed to send at given time
  - Receiver: *receiving window* – set of sequence numbers it is allowed to accept at given time
- Window size corresponds to flow control
  - May be fixed in size or adapt dynamically over time

# Simple Example: Sliding Window

- Simple sliding window example for  $n=3$ , window size fixed to 1
- Sender tracks currently unacknowledged sequence numbers
  - If maximum number of unacknowledged frames is known, this is equivalent to sending window as defined on previous slide



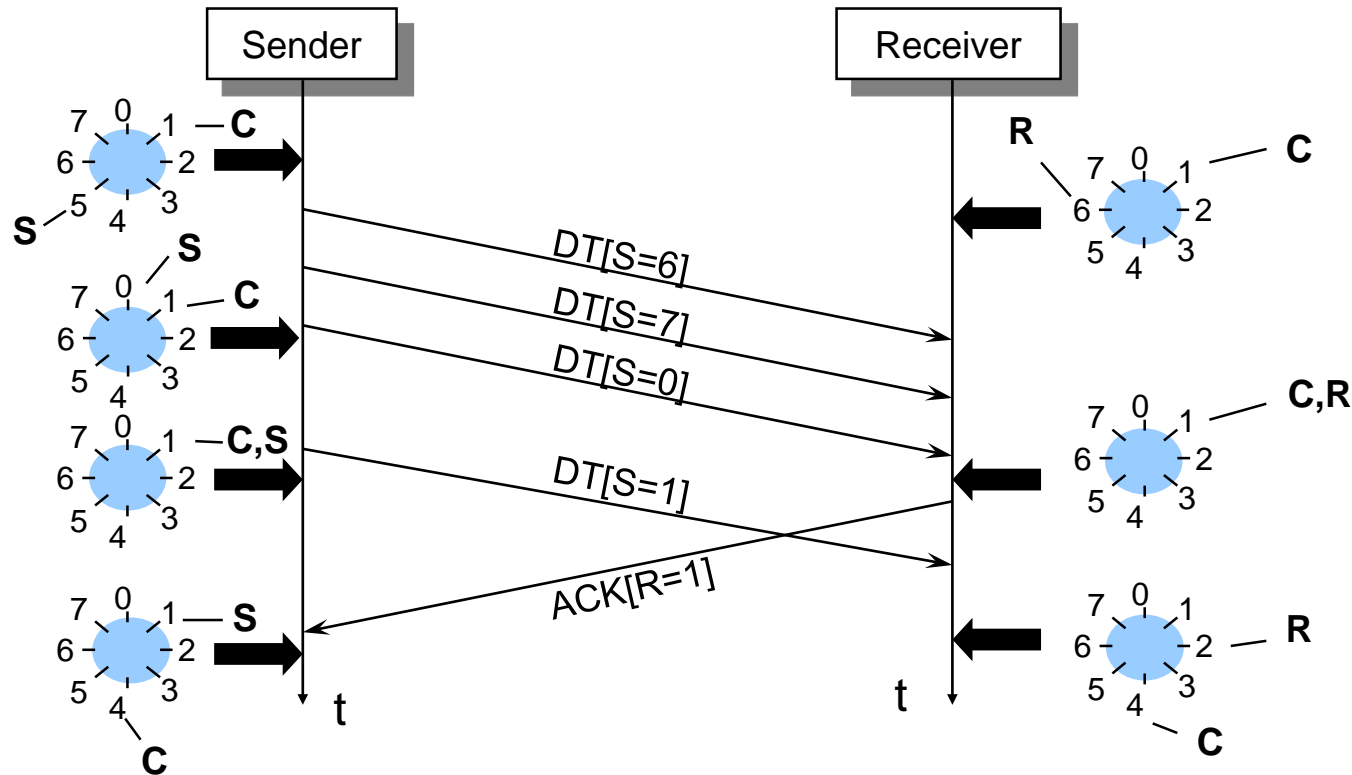
- Initially, before any frame is sent
- After first frame is sent with sequence number 0
- After first frame has been received
- After first acknowledgement has arrived





# Advanced Example: Sliding Window

- Sender credit (window size) = 4




S: Sequence number (last sent packet)

R: Next expected sequence number = Acknowledges packets up to R-1

C: Upper window limit (current max. sequence number)

➤ Disadvantage: Coupling of flow control and error control


- Assumption:
  - Link layer should deliver all frames correctly and in sequence
  - Sender is pipelining packets to increase efficiency
- What happens if packets are lost (discarded by CRC)?
- With receiver window size 1, all following packets are discarded as well!



- With receiver window size 1, all frames following a lost frame cannot be handled by receiver
  - They are out of sequence
  - They cannot be acknowledged, only ACKs for the last correctly received packet can be sent
- Sender will timeout eventually
  - All frames sent in the meantime have to be repeated
  - Go-back-N (frames)
- Quite wasteful of transmission resources
- But saves resources at receiver

# Selective Repeat

- Suppose we invest into a receiver that can buffer packets intermittently if some packets are missing
  - Corresponds to receiver window larger than 1
- Resulting behavior:



- Receiver explicitly informs sender about missing packets using *Negative Acknowledgements* (NACKs)
  - Sender selectively repeats the missing frames
  - Once missing frames arrive, they are all passed to network layer
- More resources used at receiver, less overhead in case of error

- So far, simplex operation at the (upper) service interface was assumed
  - Receiver only sent back acknowledgements, possibly using duplex operation of the lower layer service
- What happens when the upper service interface should support full-duplex operation?
  - Use two separate channels for each direction (SDMA)
    - Wasteful on bandwidth/resources
  - Interleave ACKs and data frames in a given direction (TDMA)
    - Better, but still some overhead
  - *Piggybacking*: Put ACKs from A to B into data frames from B to A (as part of B's header to A)
    - Minimal overhead
    - We'll see this principle again on layer 4 with TCP!

- Most problems in the link layer are due to errors:
  - Errors in synchronization require non-trivial framing functions
  - Errors in transmission require mechanisms to
    - Correct them so as to hide from higher layers
    - Detect them and repair them afterwards
- Flow control is often tightly integrated with error control in commonly used protocols
  - It *is* a separate function and can be implemented separately
- Connection setup/teardown still has to be addressed
  - Necessary to initialize a joint context for sender and receiver

8. Networked Computer & Internet

9. Host-to-Network I

**10. Host-to-Network II**

11. Host-to-Network III

12. Internetworking

13. Transport Layer